



Large-scale distributed computing systems

Week 3

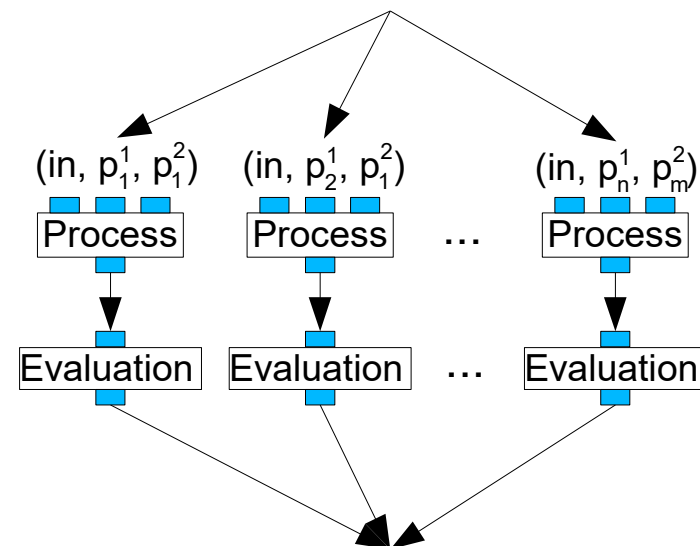


January 2017

Johan Montagnat
CNRS, I3S, SPARKS

<http://www.i3s.unice.fr/~johan/>

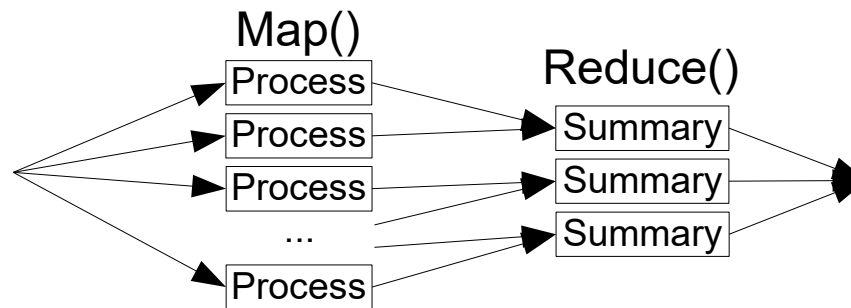
Distributed application models

- ▶ Embarrassingly parallel applications
 - ▶ Embarrassing size
 - Scalability is the challenge
 - ▶ Trivial parallelisation
- ▶ Parameters Sweep
 - ▶ Explore process parameters space
 - ▶ Many optimization problems
 - ▶ Combinatorial




Distributed application models

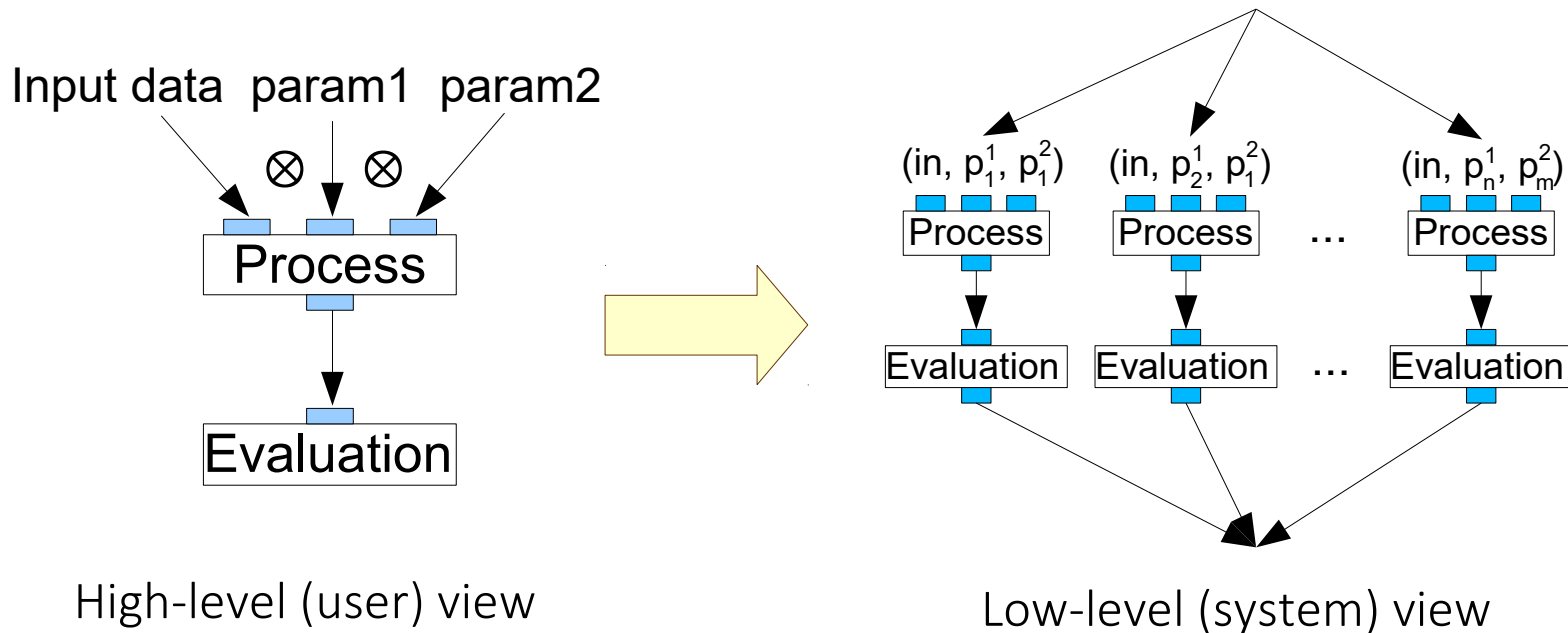
- ▶ Map-Reduce framework
 - ▶ Two steps simple parallelization framework
 - Map() function decomposes computation
 - Reduce() function combines results
 - ▶ Backed-up by fault-tolerant and scalable support tools
 - e.g. Hadoop



- ▶ SPMD (Single Processing, Multiple Data)
 - ▶ Exploit large data sets parallelism
 - ▶ Independent computations on different data items
 - ▶ Simple process: embarrassingly parallel
 - ▶ Complex process: pipelines / workflows



Workflows

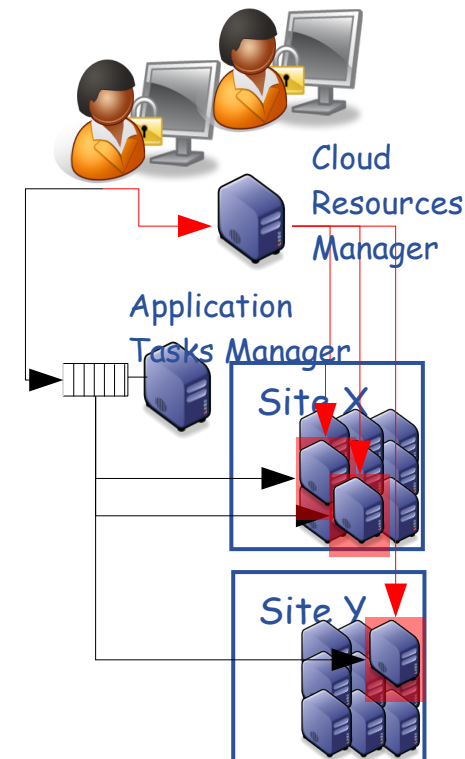


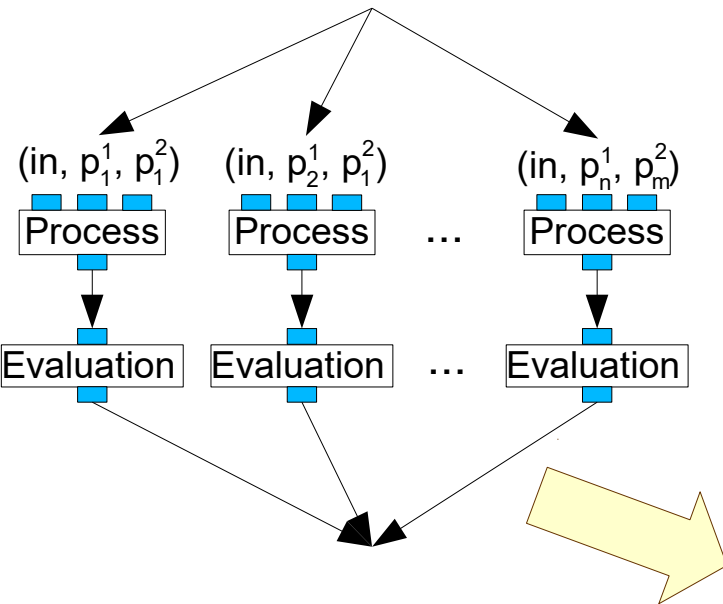


- ▶ Abstract representation for the computational process
- ▶ Close to user problem modeling view
- ▶ Compact, humanly tractable

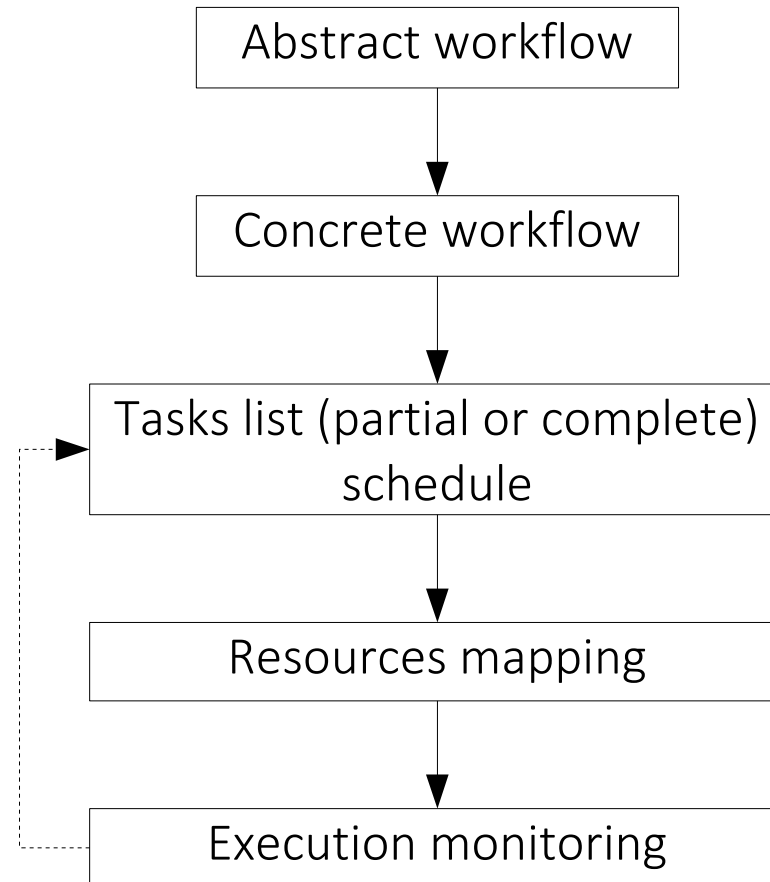
- ▶ Concrete (executable) representation
- ▶ Close to system description needs
- ▶ Detailed

Runtime

- Scheduling and mapping of workflow on computational resources



Workflow manager



- Different representations lead to different scheduling and mapping requirements


Workflow representations

Definitions of workflows

- ▶ Workflow Management Coalition:
 - ▶ "The automation of a business process, in whole or part, during which documents, information or tasks are passed from one participant to another for action, according to a set of procedural rules"
- ▶ A workflow is a **graph** representing dependencies:
 - ▶ Data and control links
 - ▶ Control structures
- ▶ A workflow links the **components** of an application:
 - ▶ Services (Web-Services, DIET services, ...)
 - ▶ Tasks (JDLs)
 - ▶ Local codes
 - ▶ Human activities
- ▶ Workflows for distributed infrastructures
 - ▶ Human scale activities, legacy codes coupling

Illustration

- ▶ Workflow




- ▶ algorithm description

```
A
then ((B then D) and (C then E))
then F
```

- ▶ Contains

- ▶ Sequential (dependent) and parallel (independent) components
- ▶ Implicit list of tasks to be computed
- ▶ Command invocation details

- ▶ Data flow



- ▶ Tasks computation


```
d = F(D(B(A(a))), E(C(A(a))))
e = F(D(B(A(b))), E(C(A(b))))
f = F(D(B(A(c))), E(C(A(c))))
```

Workflow representation

- ▶ Any programming language could be adopted...
 - ▶ From C to Makefiles
- ▶ ...but all do not fit well a distributed computing environment
 - ▶ Distributed computing workflows target human-scale activities
 - ▶ Prototyping is the rule
 - ▶ Scientific applications with heavy user communities and large data sets require high performance and/or high throughput
- ▶ Workflow languages for distributed systems
 - ▶ Coarse-grained
 - ▶ Data intensive
 - ▶ Heavy legacy code
 - ▶ Interfaced to external Job / Workload manager(s)
- ▶ Language simplicity vs expressiveness tradeoff
 - ▶ Separate design and enactment phase


Expression of parallelism

- ▶ Three levels of parallelism




Workflow parallelism
(implicit in workflow graph)

$\{D_0, D_1, D_2 \dots\}$



Data parallelism



Services parallelism
(pipelining)


- ▶ The expression of data and service parallelisms depend on the workflow language considered
- ▶ The workflow engine has to be multi-threaded (or to perform asynchronous calls)
- ▶ Data segments are puzzled

Workflow representation languages

- ▶ Directed task graphs
 - ▶ Simple dependency graphs
 - ▶ No complex structure
- ▶ Structured languages
 - ▶ Graphs with control structures
 - ▶ Scripts
- ▶ Data-driven workflows
 - ▶ Data exchanges are implicit dependencies
 - ▶ Data-driven computing model

Directed task graphs

- ▶ A task graph defines **precedence constraints** in a task set
- ▶ A task graph is a Directed Acyclic Graph (DAG)
- ▶ Ex: Condor DAGMan (<http://www.cs.wisc.edu/condor/dagman>)



Job A A.condor ①

Job B B.condor

Job C C.condor

Job D D.condor

Script PRE A top_pre.csh

Script PRE B mid_pre.perl \$JOB

Script POST B mid_post.perl \$JOB \$RETURN

Script PRE C mid_pre.perl \$JOB

Script POST C mid_post.perl \$JOB \$RETURN

Script PRE D bot_pre.csh

PARENT A CHILD B C ②

PARENT B C CHILD D ③


Retry C 3

Structured workflow graphs

- ▶ Control-centric, aka Business Workflow
- ▶ Example: Business Process Execution Language (BPEL)
 - ▶ From IBM (WSFL) and Microsoft (XLANG)
 - ▶ Build on Web-Services standards
- ▶ Defining the behavior of a process with a formal description of the messages exchanged by the Web-Services
- ▶ Specifies the behavior of all “partners” independently from their implementation
- ▶ The resulting process is itself a Web-Service

Software components and services


- ▶ Software components



- ▶ Define and publish a standard interface
 - ▶ Interact by message exchanges
 - ▶ Ease (dynamic) composition
 - ▶ Modularity (no compilation time dependency)
 - ▶ High level message exchange protocols (format, types...)
- ▶ Services
 - ▶ Independent, self-sufficient software components
 - ▶ Can be invoked remotely
 - ▶ Can be dynamically created and destroyed

Service-Oriented-Architectures (SOA)

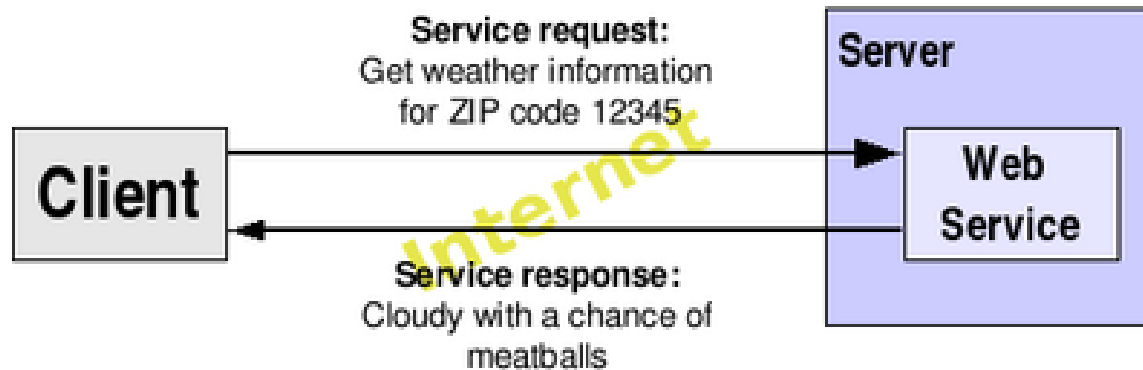
- ▶ Basic principles



- ▶ A service is an exposed piece of functionality with 3 properties:
 - (1) The interface contract to the service is **platform-independent**
 - (2) The service can be **dynamically located** and invoked
 - (3) Services maintain a relationship that minimizes dependencies (**loosely coupling**)

Services in a SOA


- ▶ Client/Server interaction



- ▶ Platform and language independent
 - ▶ Client and server programs can be written in different languages, and run in different environments
- ▶ Self-describing
 - ▶ Only location is needed to invoke a service
 - ▶ Loosely coupled
- ▶ Based on the adoption of common standards

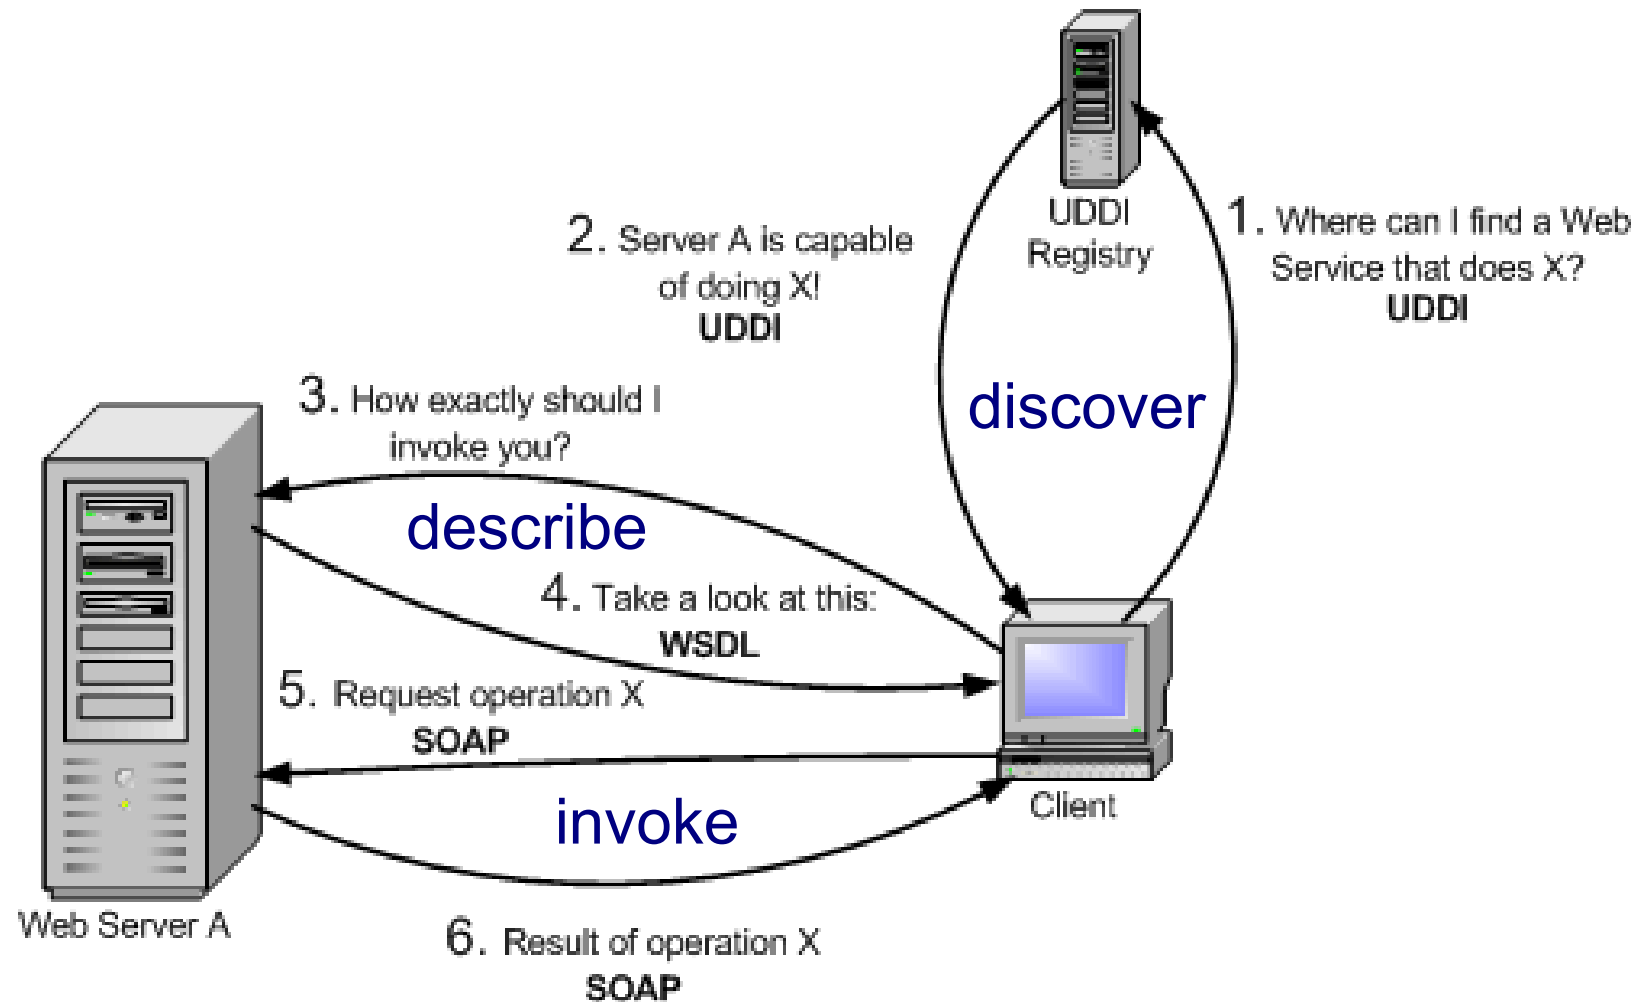
Web-Services standard

- ▶ Standardized by the W3C:
 - ▶ Contract / Interface format
Web **S**ervices **D**escription **L**anguage
 - ▶ Messages format
Simple **O**bject **A**ccess **P**rotocol
 - ▶ Discovery format
Universal **D**escription **D**iscovery & **I**ntegration
- ▶ Based on XML
 - ▶ Text format
 - ▶ Platform/language independence


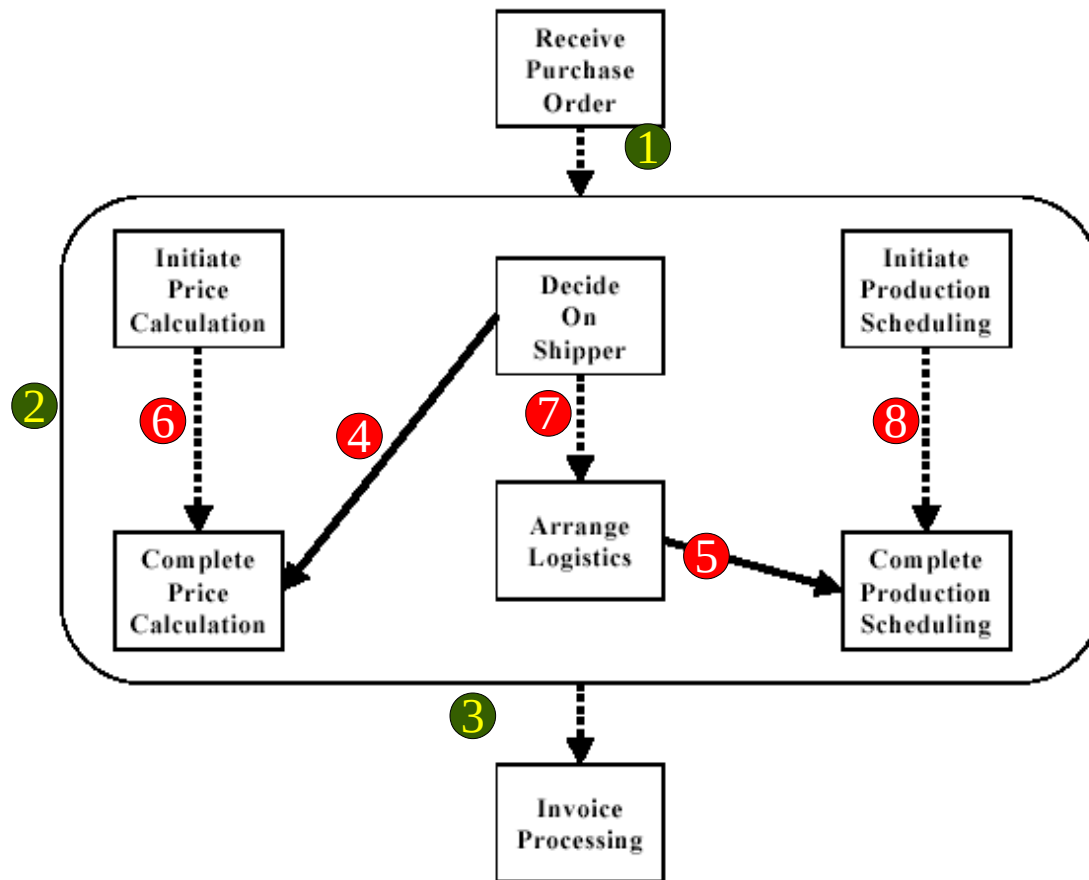


Web services

- Typical use case



BPEL web services workflow example



declarations

activities

```

<process name="purchaseOrderProcess">
  <partnerLinks>
  </partnerLinks>
  <variables>
  </variables>
  <faultHandlers>
  </faultHandlers>
  <sequence>
    <receive> 1
    </receive>
    <flow>
      <links> 4 5
      </links>
      <sequence> 6
      </sequence>
      <sequence> 7
      </sequence>
      <sequence> 8
      </sequence>
    </flow>
    <reply> 3
    </reply>
  </sequence>
</process>

```

BPEL: structured activities

- ▶ Can contain other activities:

<sequence>

sequential execution

<flow>

parallel execution

<pick>

blocks the execution until a message/timeout occurs

<switch>

selects an activity according to a condition

<while>

iteration

<scope>

defines an activity with its own variables, handlers, ...

BPEL: communicating activities

- ▶ Interact with the partners of the workflow:

<invoke>

sends a message to a port of a partner

<receive>

blocking wait of a message

<reply>

sends a message replying to a received message (by <receive>)

BPEL: other activities

- ▶ Other activities

- <assign>

- assigns a value to a variable*

- <wait>

- waits for a given duration or until an instant*

- <terminate>

- terminates the process*

- <compensate>

- executes the compensate field (called by a fault handler)*

- <throw>

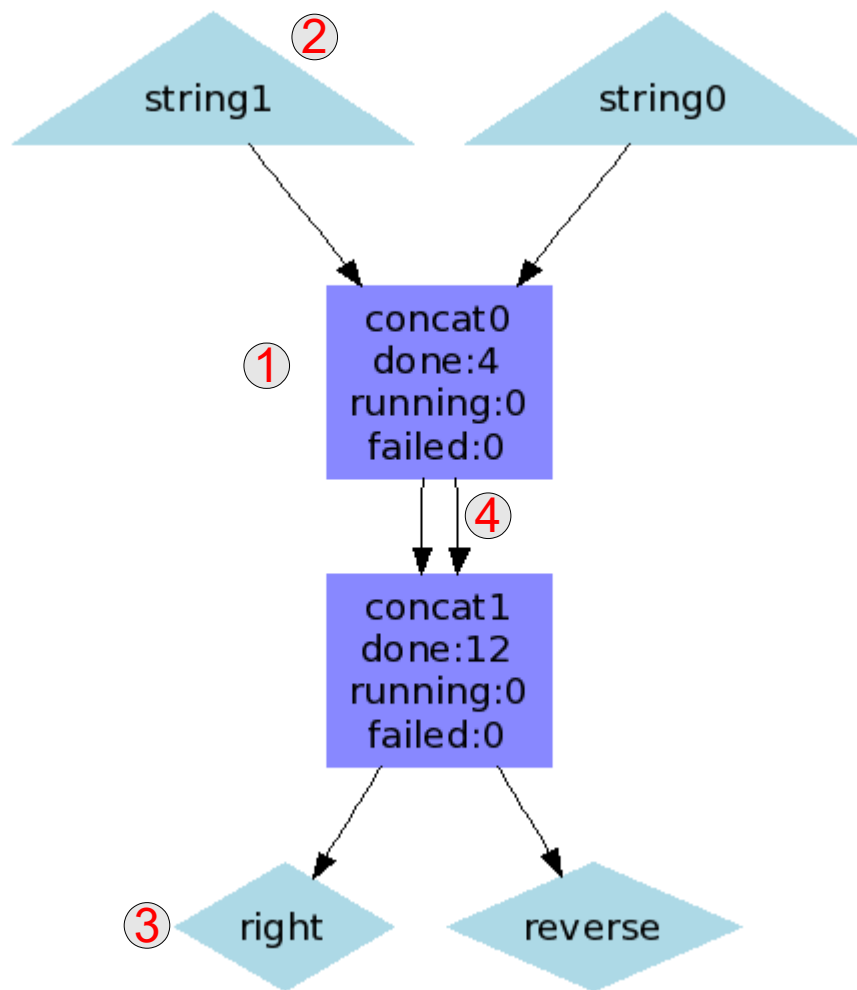
- throws an exception*

- <empty>

- nop*

Data-driven workflows

- ▶ Data-centric, aka scientific workflows
- ▶ Example: Simple Concept Unified Flow Language (Scufl)



```

<processor name="concat0"> 1
  <arbitrarywsdl>
    <wsdl>http://colors.unice.fr/concat_service.wsdl</wsdl>
    <operation>concat</operation>
  </arbitrarywsdl>
  <iterationstrategy>
    <cross>
      <iterator name="string2" />
      <iterator name="string1" />
    </cross>
  </iterationstrategy>
</processor>

<processor name="concat1"> ... </processor>

<source name="string1" /> 2

<sink name="right" /> 3

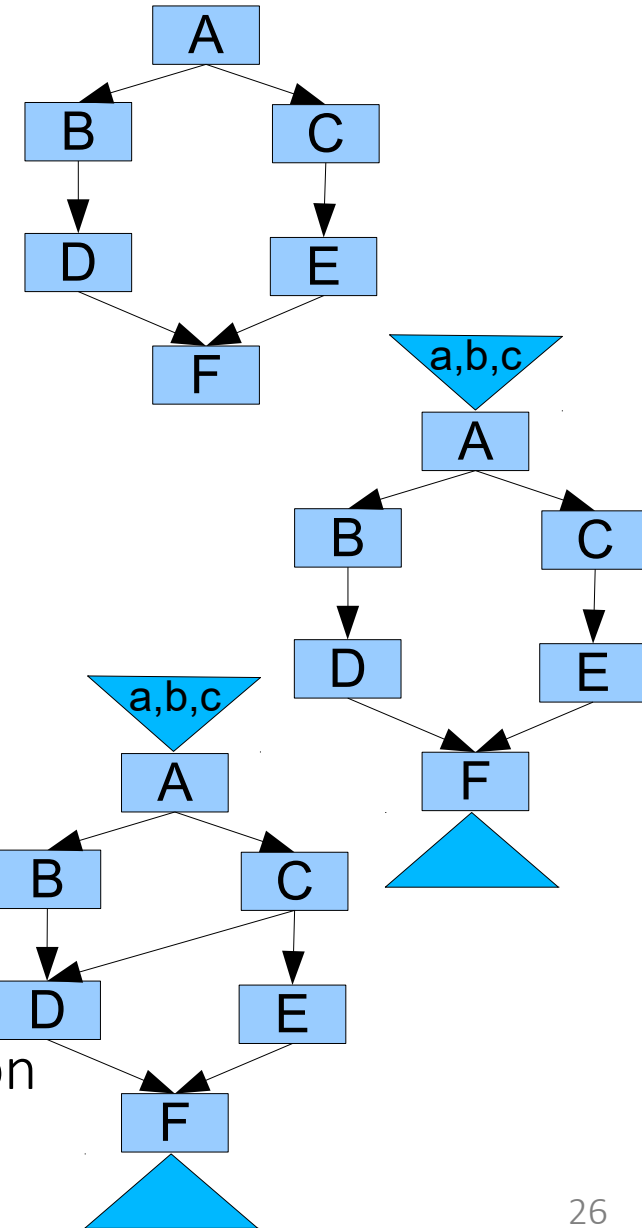
<link source="concat0:reverse" sink="concat1:string1" /> 4
  
```


Taxonomy of workflow approaches

- ▶ Workflow descriptions
 - ▶ Business workflows languages (*e.g.* BPEL) → Control-centric
 - ▶ Scientific workflows languages (*e.g.* Scufi) → Data-centric
- ▶ Scientific workflows approaches
 - ▶ **Service-based** workflows (*e.g.* Scufi)
 - Independent expression of processings and input data
 - Dynamic description
 - Simple representations, complex optimization
 - ▶ **Task-based** workflows (*e.g.* DAGMan)
 - Tasks mixes data and processings
 - Static description
 - Complex representations, simple optimization

Scripting approach

- ▶ Any programming language may be use to describe a workflow
 - ▶ Workflow representation buried in code
- ▶ May use parallel control structures
 - ▶ Explicit description of parallel tasks
 - `Dopar(Activity(D0), Activity(D1))`
 - ▶ Parallel loops
 - `D = { D0, D1, D2 }`
`ParallelForeach(d in D) {`
 `Activity(d)`
`}`
 - ▶ Parallel threads
 - `Fork() / Join()`
- ▶ “1D” (linear) code representation
 - ▶ Opposed to 2D graphs
 - ▶ Parallelism implicitly expressed in 2nd dimension



Future variables

- ▶ Traditional languages have blocking assignment instructions
 - ▶ `a = f(0);` // variable assignment is blocking:
 `b = g(0);` // execution of `f()` completes and `a` is
 // assigned before `g()` is executed
 `c = h(a, b);` // `h` executed once `a` and `b` values are known
- ▶ Future variables are non-blocking assignment variables
 - ▶ Value read is blocking though
 - To preserve computations coherency
 - ▶ `future a, b;`
 `a = f(0);` // variable assignment is non-blocking:
 // `f(0)` is evaluated asynchronously
 `b = g(0);` // `g(0)` is evaluated immediately
 `c = h(a, b);` // variable read is blocking
 // `a` and `b` are evaluated before `h(a, b)`


Future variables

- ▶ Equivalent representations

- ▶ `Dopar { a = f(0), b = g(0) };
c = h(a, b);`

- ▶ `f1 = Fork(a = f(0));
f2 = Fork(b = g(0));
join(main, f1, f2);
c = h(a, b);`

- ▶ Graph-based



- ▶ Future variables introduce asynchronism in synchronous languages

- ▶ Implicit representation of parallelism
 - ▶ e.g. SwiftScript workflow language

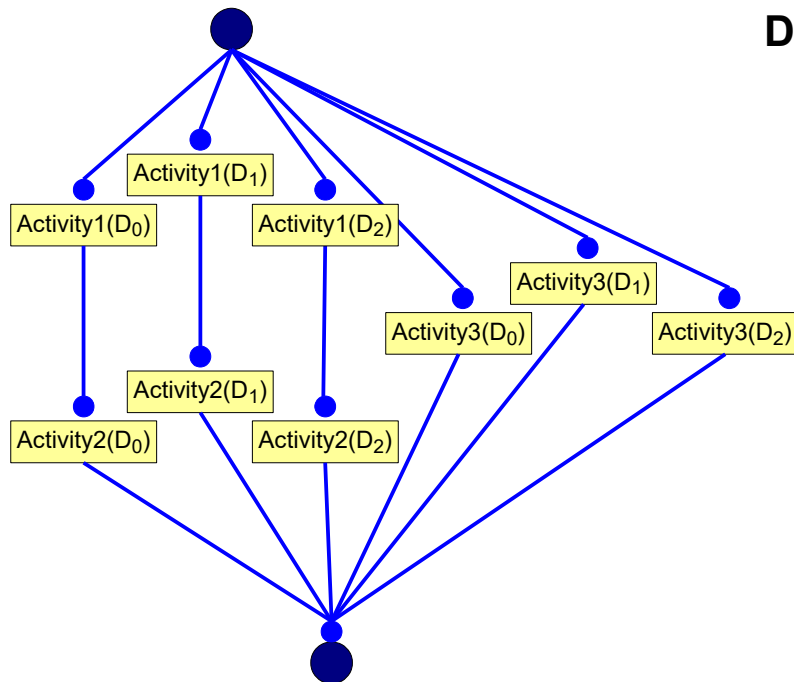
Parallel languages vs data-driven flows

- ▶ Parallel languages
 - ▶ Bounded: `foreach d in {D0, D1, D2}...`
 - ▶ Unbounded: `foreach d in D...`
 - ▶ Involves data synchronization at each data parallel service (no pipelining)
- ▶ Data-driven flows
 - ▶ Independent definition of data flows and data sets
 - ▶ Enable completely asynchronous enactment (data parallelism + pipelining): no data synchronization
 - ▶ May require explicit data synchronization barriers when this is needed
- ▶ Futures
 - ▶ Non-blocking assignment operations, blocking read
 - ▶ Data-centric approach, asynchronous execution

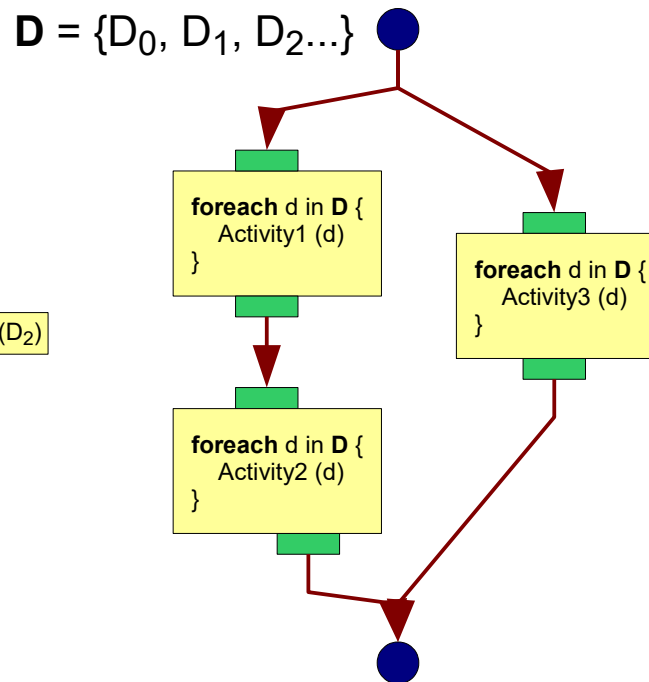
Control- vs data-driven language

- ▶ Different representations, similar semantics
 - ▶ Trade-off between workflow representation expressiveness and scheduling complexity


control ←————→ data



DAGMan: control-centric




SwiftScript: control & futures



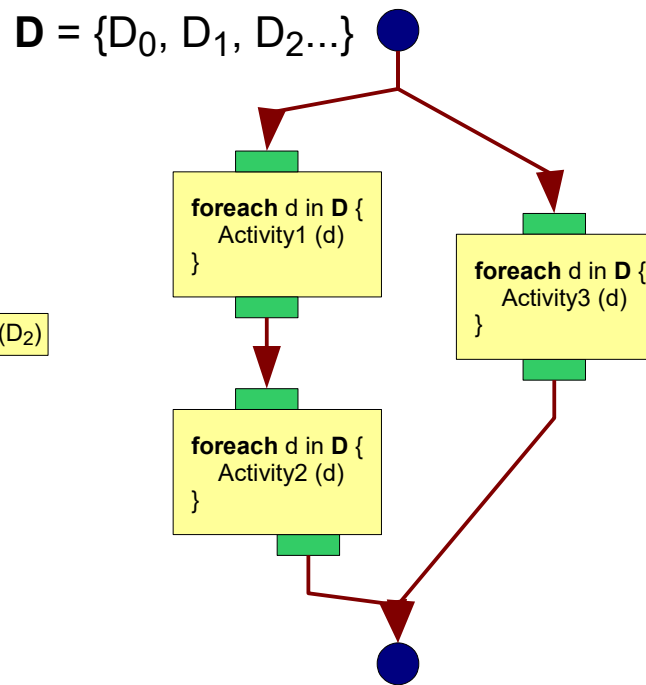
Scufl: data driven

Representation of data parallelism


- ▶ Most languages express data parallelism



DAGMan: control-centric



SwiftScript: control & futures



Scufl: data driven

- ▶ DAGs: implicit in the (large-scale) workflow graph
- ▶ Parallel languages: explicit control structures
- ▶ Data-driven: transparent

Data-driven fits data-intensive well

- ▶ Main source of parallelism comes from data sets
 - ▶ Embarrassingly parallel
 - ▶ Parameter sweep
 - ▶ Single Process, Multiple Data (SPMD)
 - e.g. Map-Reduce
- ▶ Clear separation of application logic and experiment
 - ▶ Experiment = raw data and parameter data
 - ▶ This is a clear difference with BPEL (orchestration of Web Services) or other imperative programming inspired approaches
- ▶ Implicit description of parallelism

Arrays, activities and ports

- ▶ Array programming principles

- ▶ Lightweight syntax for handling arrays:

```
X+Y  $\equiv$  foreach i in indices(X) do  
    Xi + Yi  
done
```

- ▶ Convenient representation for vectorial processors
 - ▶ Extendable to any operation on arrays of values
 - ▶ Use nested arrays **x={{1, 2},{-1,-2}}** is a 2-nesting levels array

- ▶ Activities with array parameters

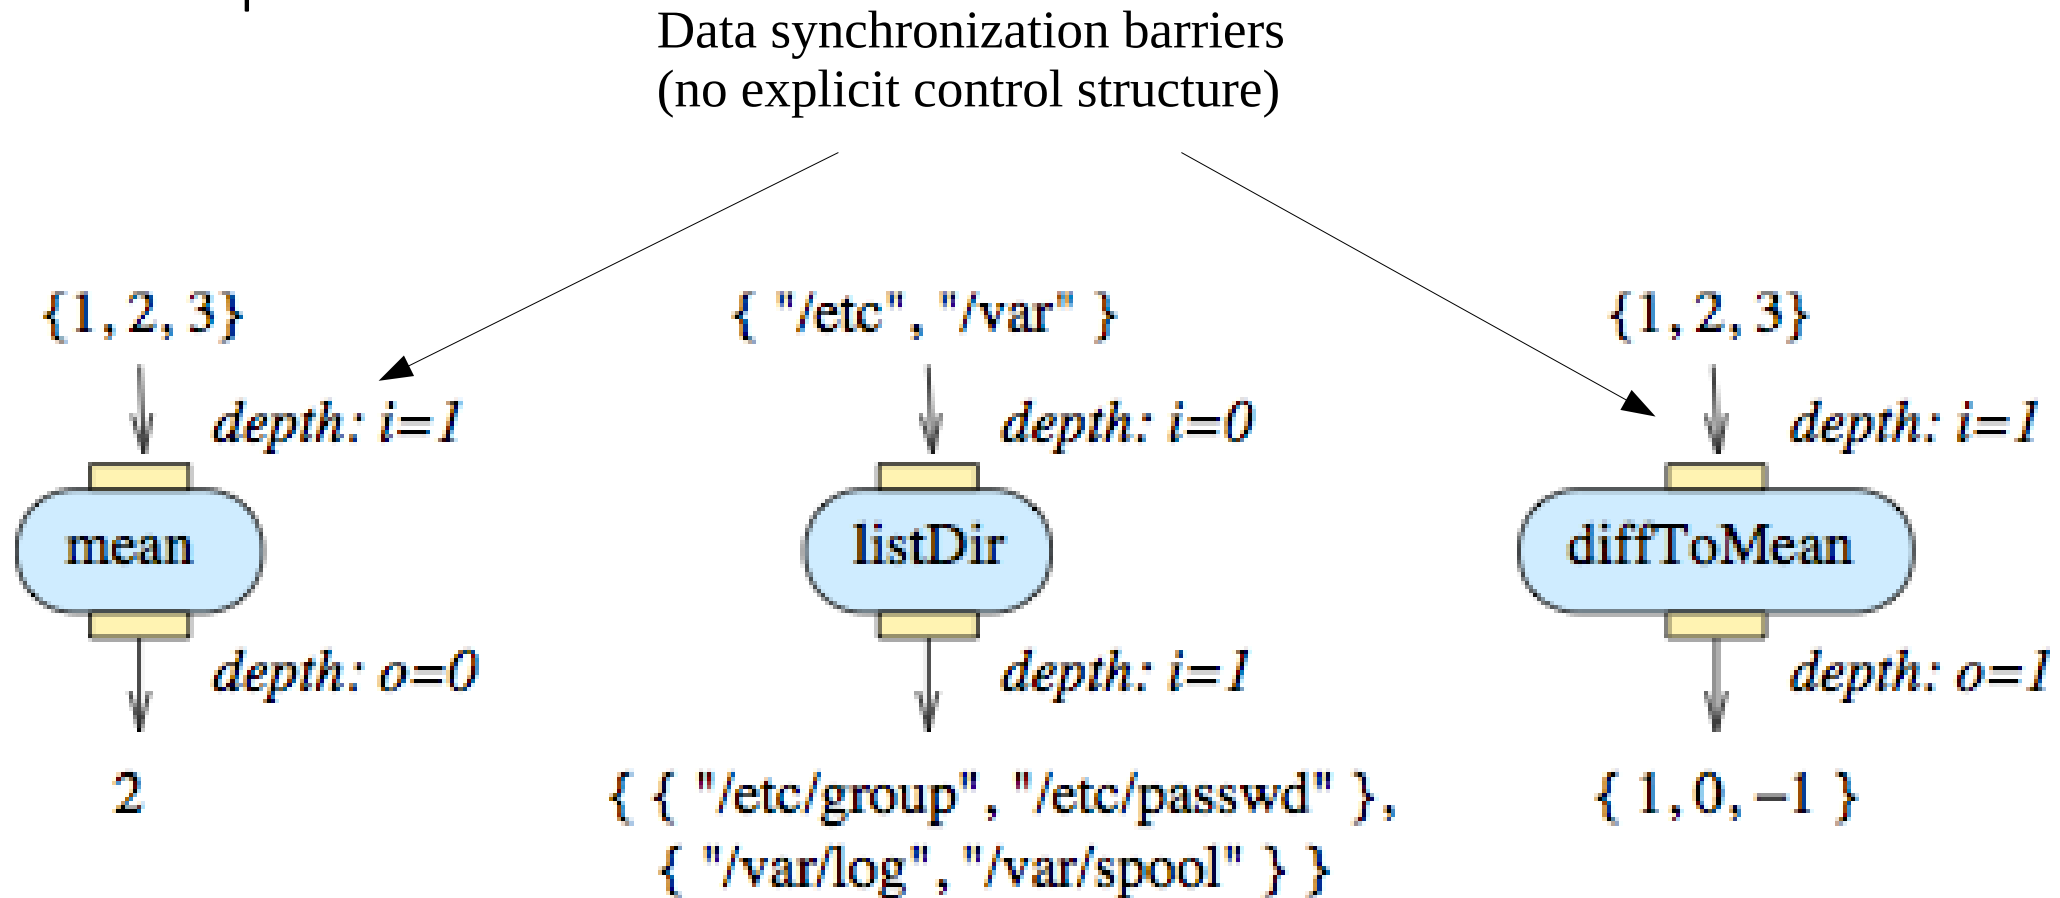
- ▶ **A(X)** \equiv { **A(X0)**, ..., **A(Xn)** }
 - ▶ **A({{1, 2},{-1,-2}})** \equiv { {**A(1)**, **A(2)**}, {**A(-1)**, **A(-2)**} }

- ▶ Ports depth


- ▶ Activity ports have a defined depth that corresponds to the number of array nesting levels being synchronized:
 - ▶ **Mean({1, 2, 3}) = 2** if **Mean** has input port depth 1

Nested arrays and ports depth


- Input arrays nesting levels are combined with activities ports depth



Iteration strategies



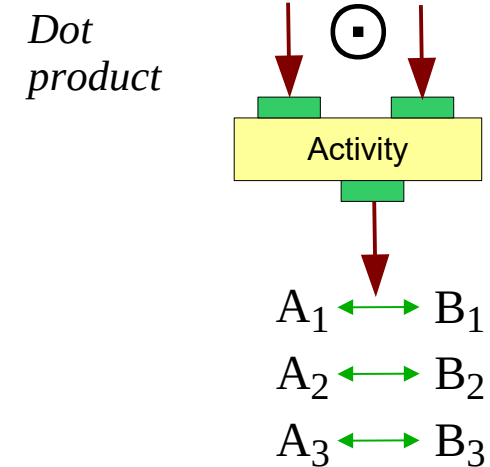
$\{A_1, A_2, A_3\}$ $\{B_1, B_2, B_3\}$



Cross product

$A \otimes B$

$\{A_1, A_2, A_3\}$ $\{B_1, B_2, B_3\}$



Dot product

$A \odot B$

► In Scufi

► Parallel language


```
foreach a in A
foreach b in B
  fire Activity(a,b)
```

```
foreach i in indices(A)
  fire Activity(Ai,Bi)
```

Iteration strategies in a parallel WF

- ▶ Dot products assume ordered data sets

Dot product

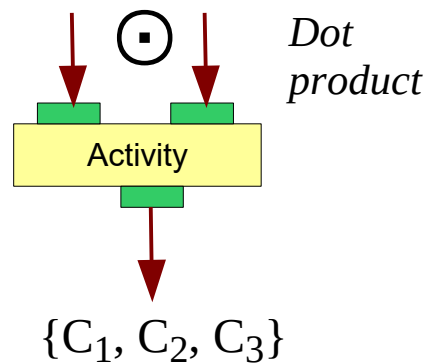


- ▶ No problem if:
 - ▶ Data parallelism is not present (order is preserved)
 - ▶ Service parallelism is not present (reordering is possible)
- ▶ Dot products in a data+service parallel execution:
 - ▶ Keep track of the data graph
 - ▶ Defines the dot product from the data graph


Iteration strategies and depths

- ▶ Iterations strategies define output indexing strategies
- ▶ They may change data nesting levels

$\{A_1, A_2, A_3\}$ $\{B_1, B_2, B_3\}$

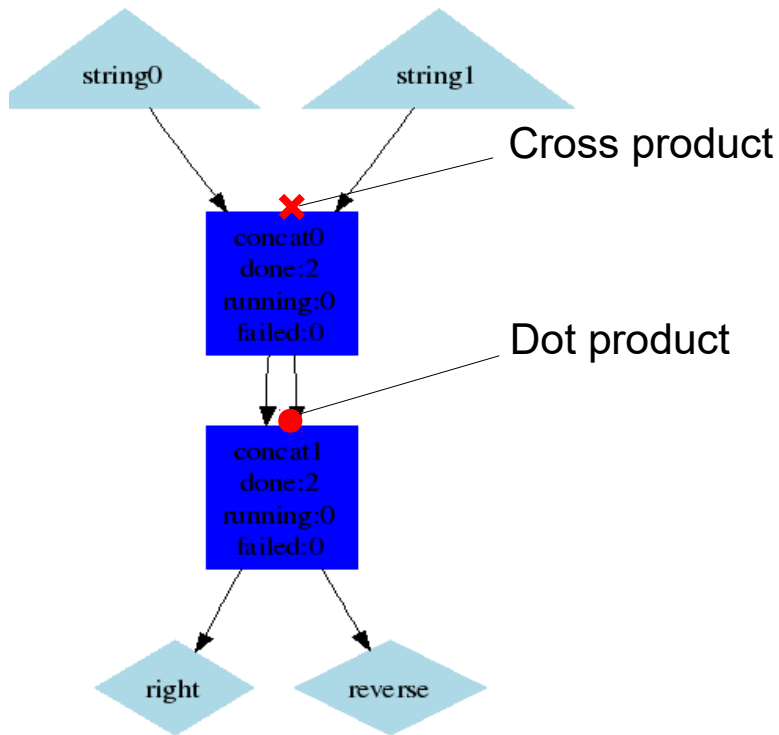


$\{A_1, A_2, A_3\}$ $\{B_1, B_2, B_3\}$

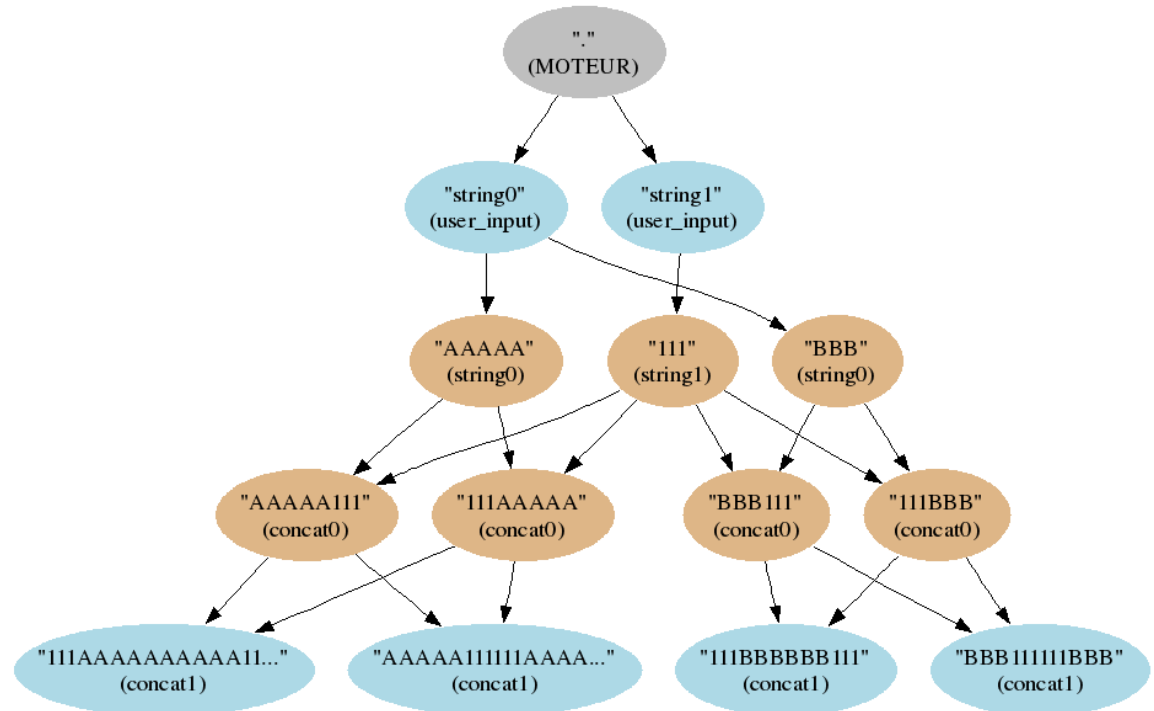


Data handling inside workflows

Services workflow




Data provenance graph




- ▶ This data representation allows to:
 - ▶ Handle dot products iteration strategies if data segments are puzzled
 - ▶ Retrieve results provenance

Data-driven workflows VS task graphs

Service-based workflow

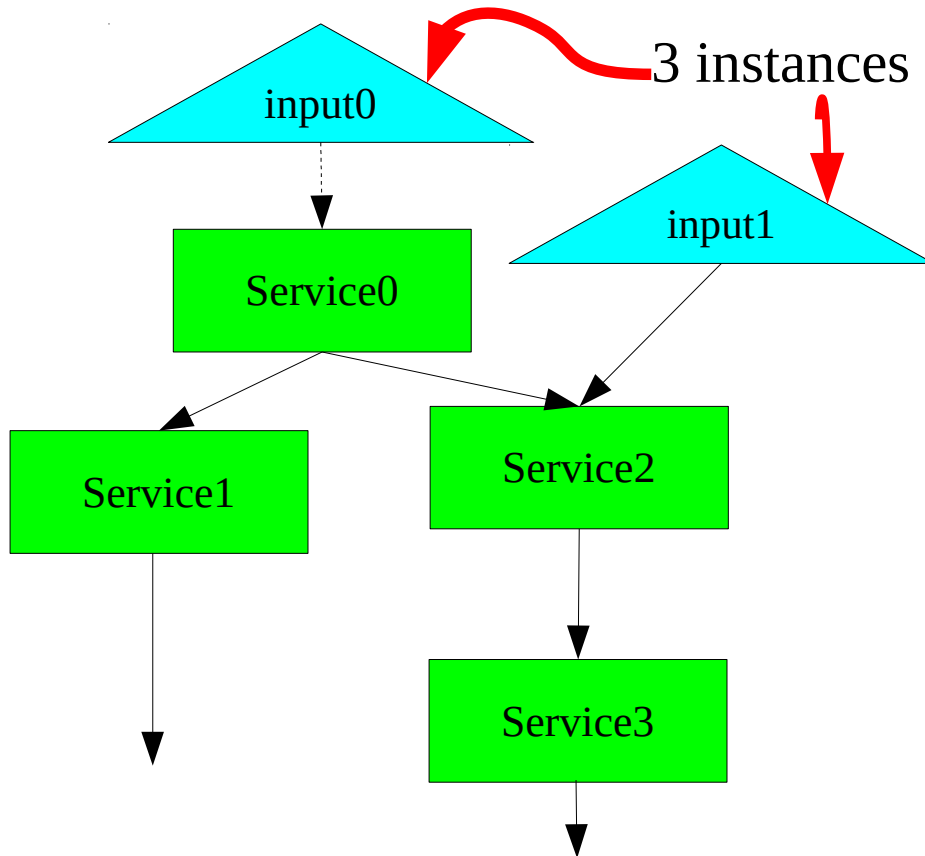


Task-graph




Data-driven workflows VS task graphs

Service-based workflow



Task-graph



- Data-driven workflows offer a data independent representation

Control structures or not?

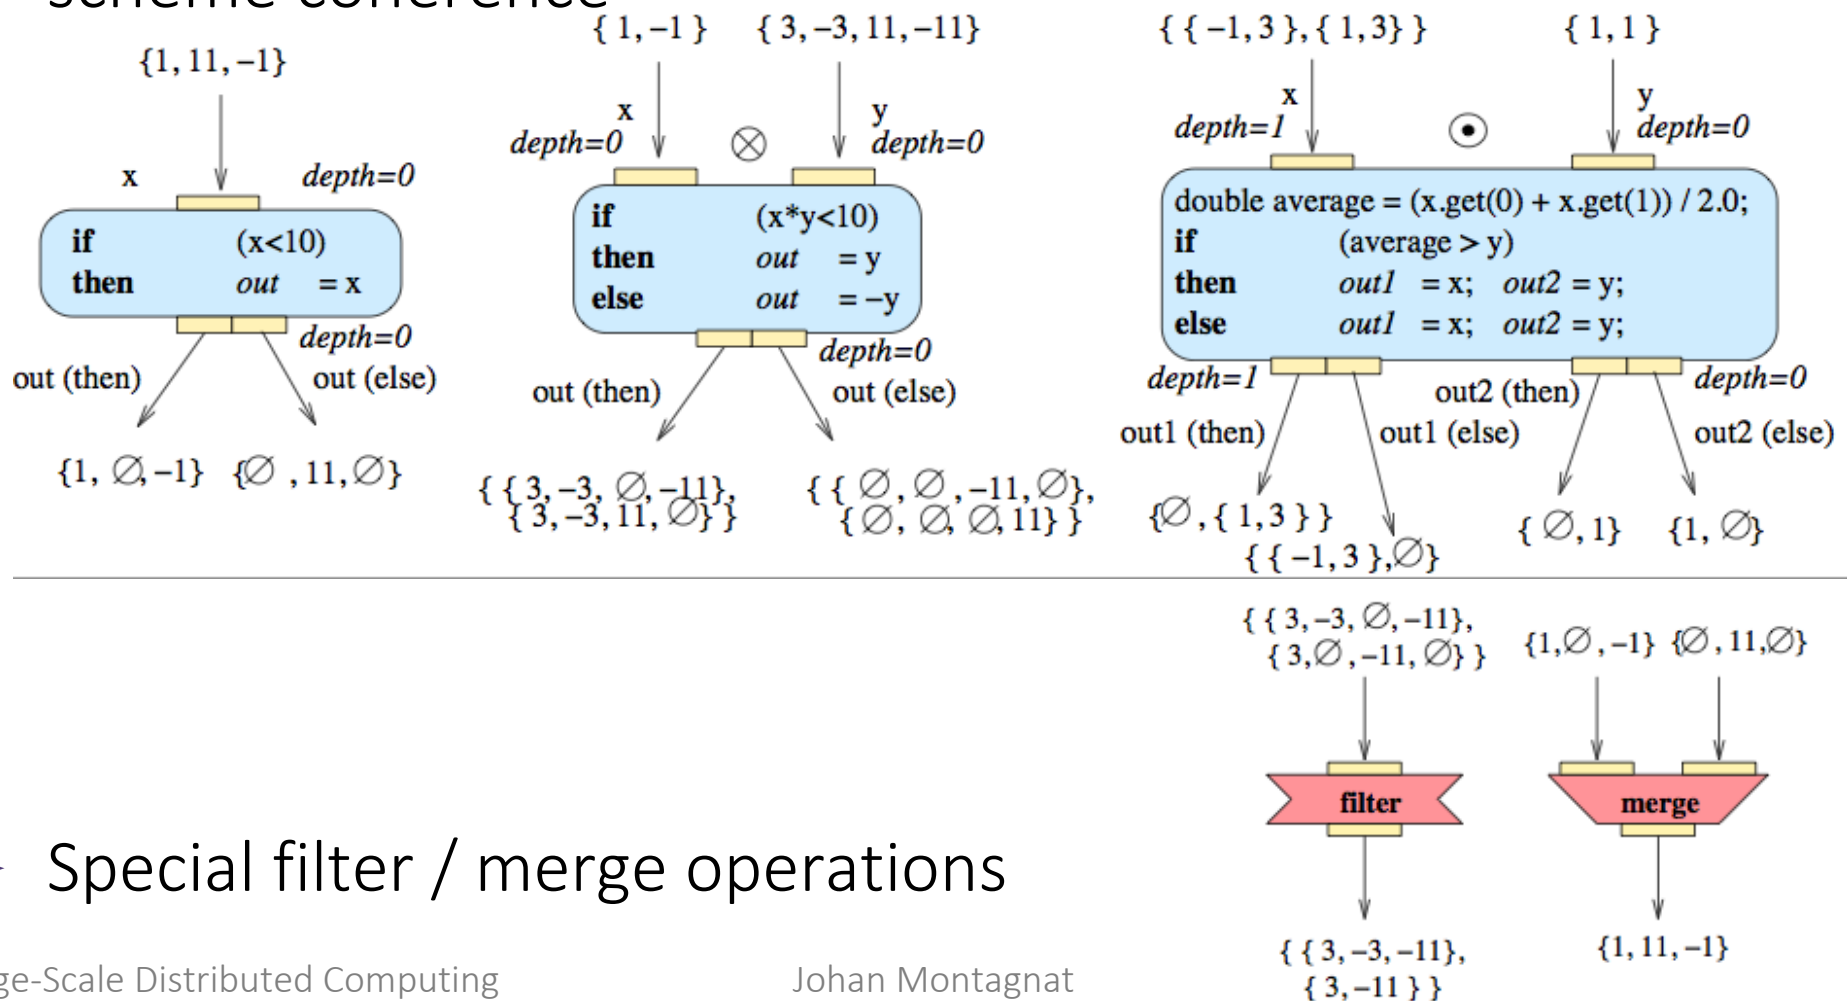
- ▶ Many data-driven / visual programming approaches do not define control structures
 - ▶ Hardly any, yet some (fail-if-false, fail-if-true)
- ▶ Directed Task Graphs do not include any loop
 - ▶ Because complex application logic is described inside the workflow activities
- ▶ Yet control on data flows is sometimes needed
 - ▶ Different execution conditions
 - ▶ Exceptions / Retry on errors at the application level
 - ▶ ...

Empty data element

- ▶ Special *void* value \emptyset
- ▶ $\mathbf{A}(\emptyset) \equiv \emptyset$
 - ▶ No evaluation of the activity (\emptyset is only known from the workflow engine)
- ▶ Iteration strategies semantics
 - ▶ $\mathbf{x} \odot \emptyset \equiv \emptyset \odot \mathbf{x} \equiv \emptyset$
 - ▶ $\mathbf{x} \otimes \emptyset \equiv \emptyset \otimes \mathbf{x} \equiv \emptyset$
- ▶ \emptyset has an index in the array it belongs to

Conditionals

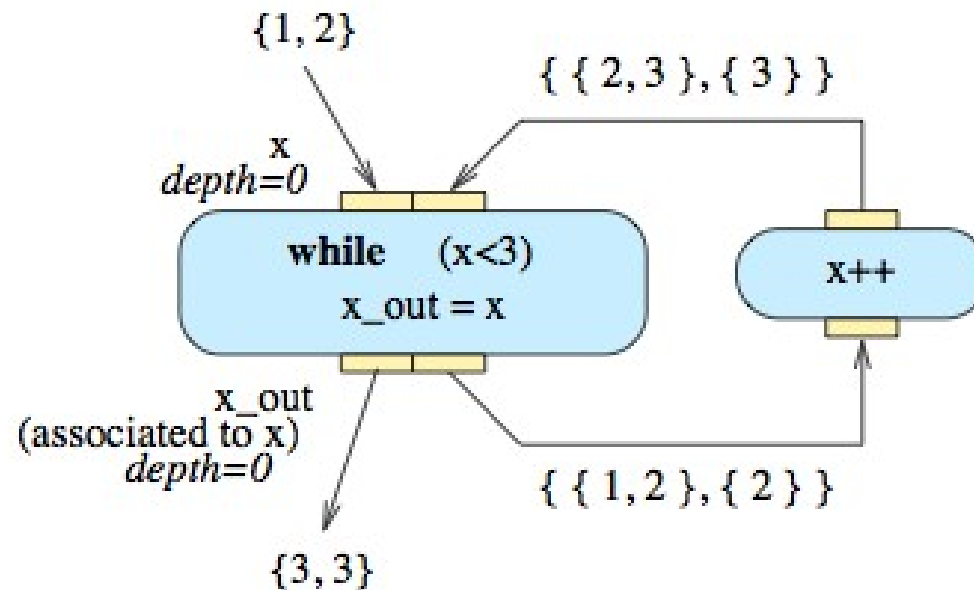
- ▶ Test expression variables mapped to input ports
- ▶ Dual “then / else” output ports
- ▶ Use of the special “void” data element to preserve indexing scheme coherence



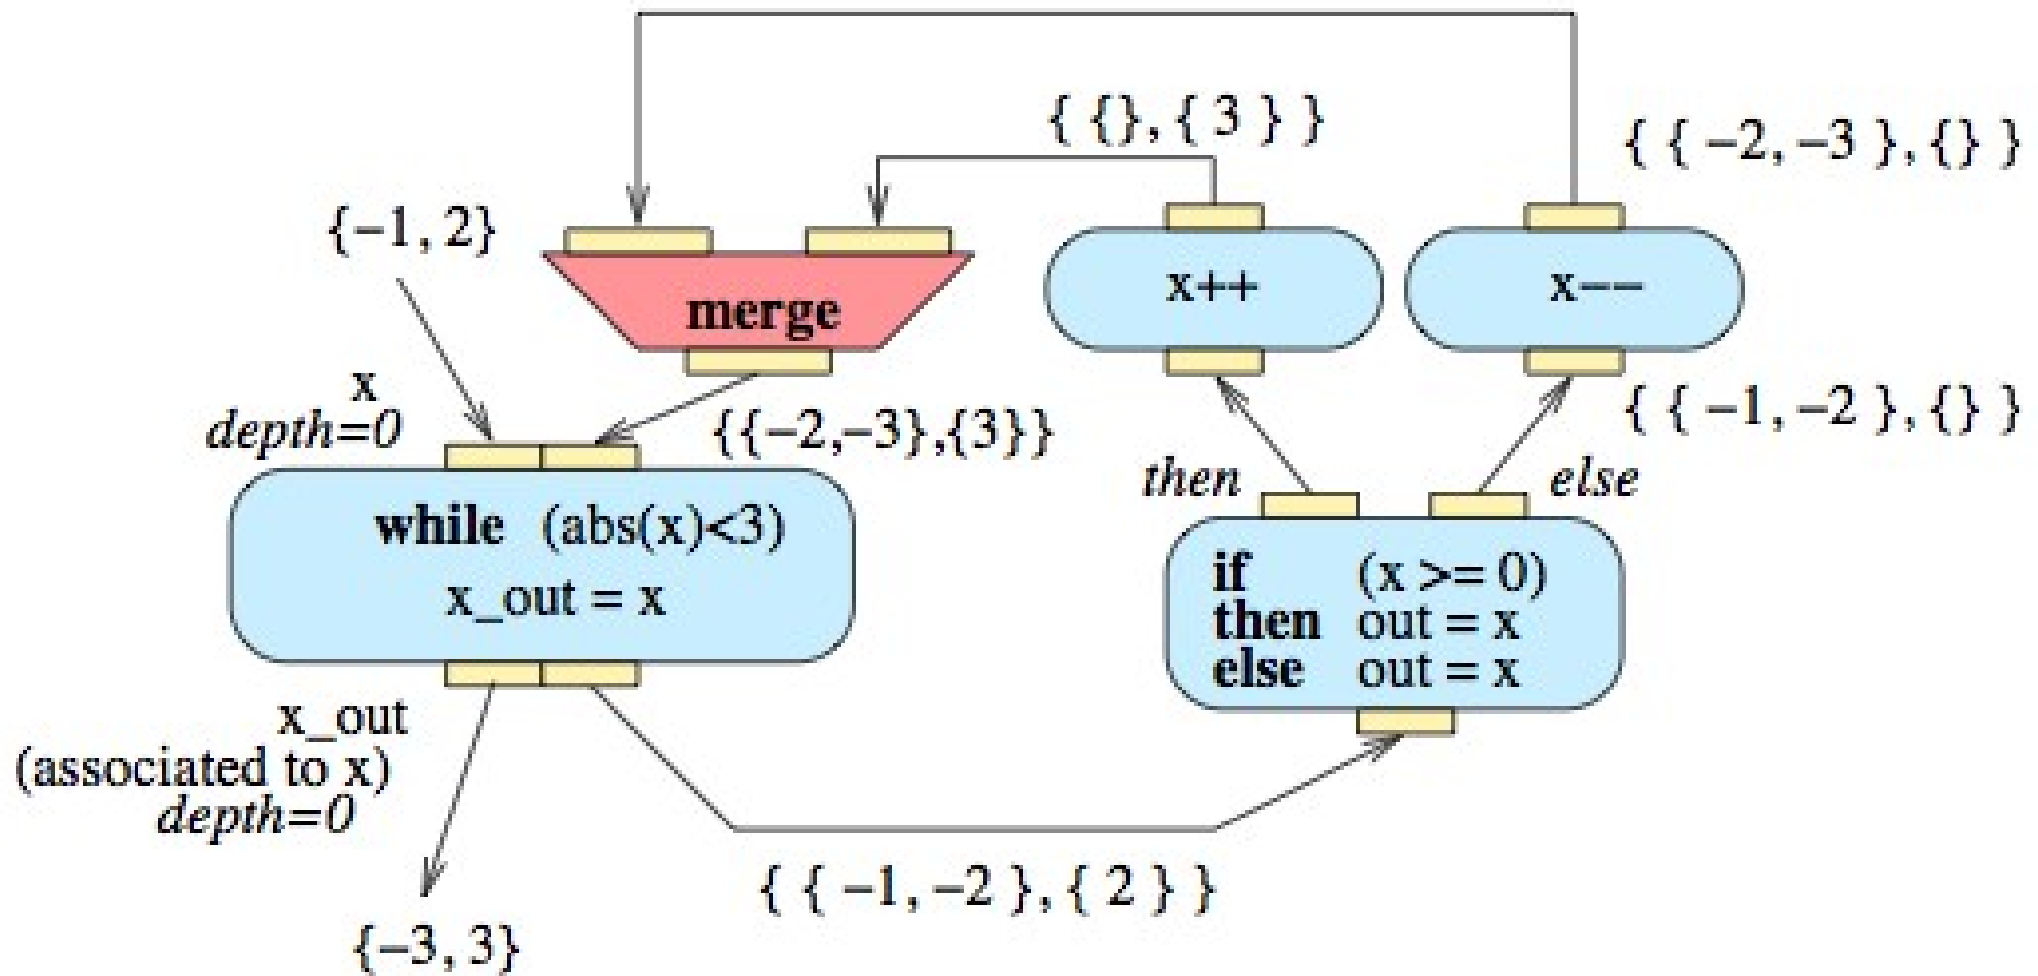
- ▶ Special filter / merge operations

Loops

- ▶ Test expression variables mapped to input ports
- ▶ Dual “inner / outer” ports
 - ▶ Input ports: initialization / iterated values
 - ▶ Output ports: iterated values / end of loop




Example




Data-driven workflows VS task graphs

- **Dynamic** data sets can be handled by service workflows:

- Loops



Conditional data set size

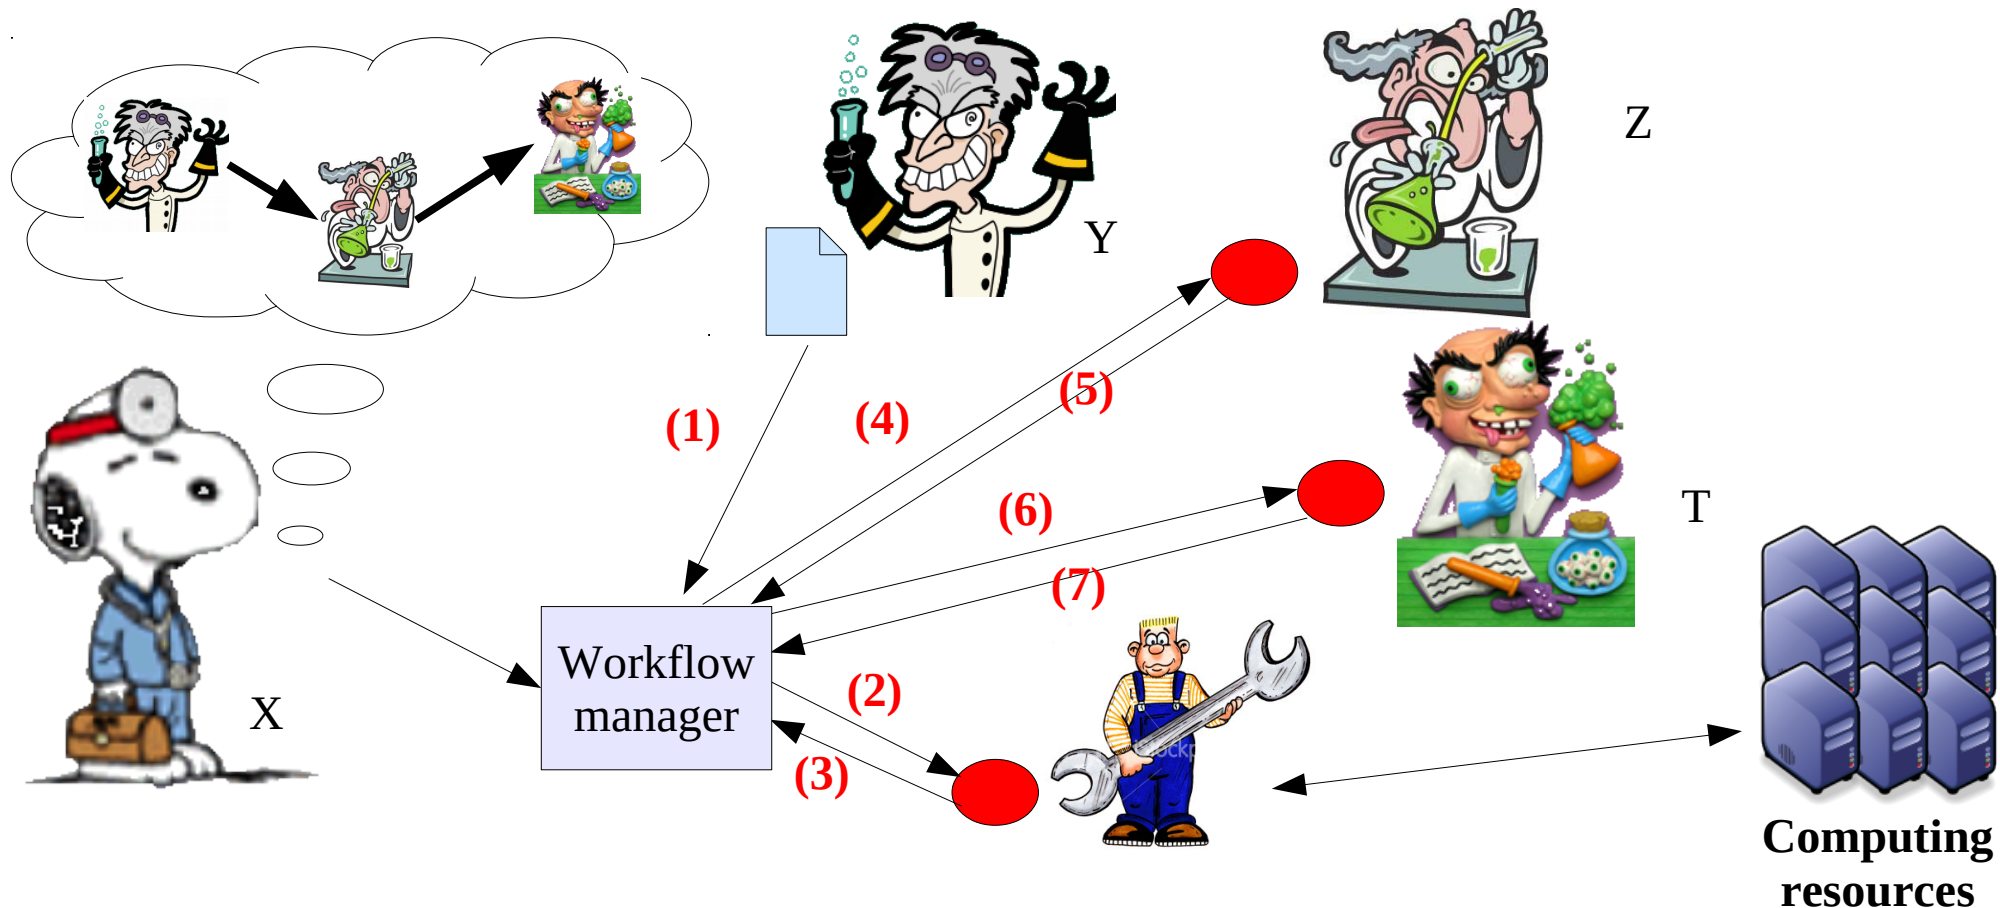


- Directed task graphs only allow **static** descriptions

Runtime

The workflow manager

- ▶ The workflow manager is a centralized engine that performs the calls to the services to execute the workflow
- ▶ It is a **generic client** to the services

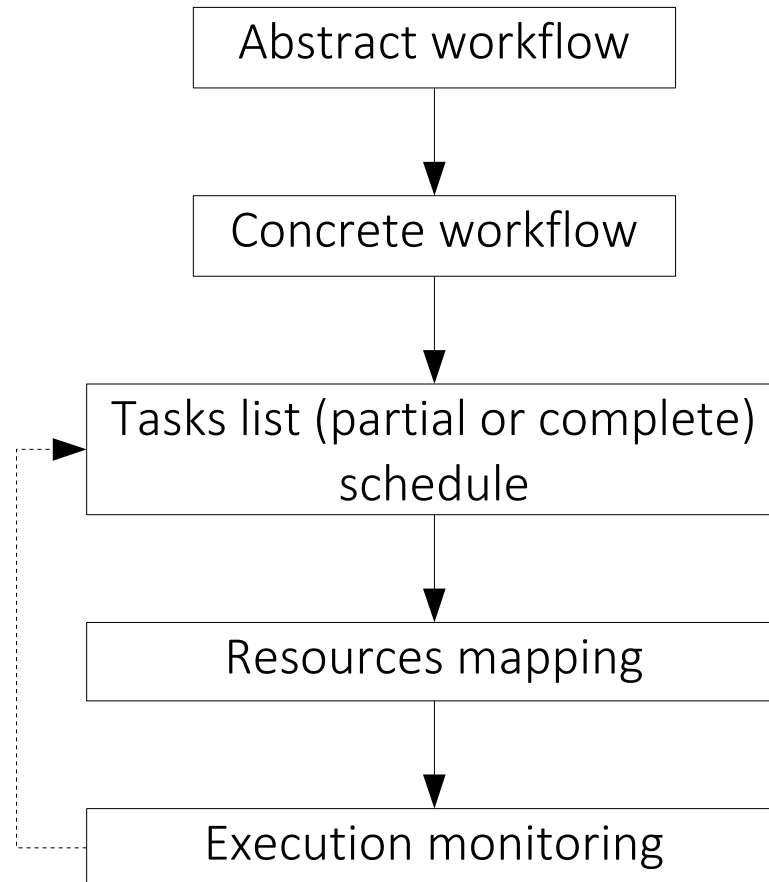


- ▶ Handling **references** to data is critical

Workflow managers

- ▶ Business workflows
 - ActiveBPEL engine, Apache ODE, ... (30 more)
 - JOpera (<http://www.jopera.ethz.ch>)
- ▶ Scientific workflows
 - ▶ Service Based
 - Taverna (<http://taverna.sourceforge.net>)
 - Triana (<http://www.trianacode.org/>)
 - Kepler (<http://kepler-project.org/>)
 - MOTEUR (<http://modalis.polytech.unice.fr/softwarewares/moteur>)
 - P-Grade portal (<http://portal.p-grade.hu>)
 - ▶ Task graphs
 - DAGMan (Pegasus and Chimera on top of it)
 - DIET MA-DAG

Workflow manager






- Different representations lead to different scheduling and mapping requirements

Scheduling


The scheduling problem

- ▶ Scheduling aims at finding a task execution and resources allocation planning to optimize one criterion or more.

- ▶ Example:

- ▶ 3 tasks:   
- ▶ 2 resources: R_1 and R_2
- ▶ Criterion: makespan (total execution time)

- ▶ Gantt charts of two schedules



Additional constraints

- ▶ Precedence constraints between tasks
 - ▶ Workflows/DAGs
- ▶ Communication costs
 - ▶ Data transfers between tasks
- ▶ Heterogeneous resources
 - ▶ Machines performances (CPU/memory)
 - ▶ Network bandwidth
- ▶ Dynamicity
 - ▶ Resources creation/deletion
 - ▶ Tasks creation/deletion

Scheduling task graphs

[Legrand, Robert 2003]


- ▶ Without communication costs:
 - ▶ The problem without resources limitations is polynomial
 - Scheduling ASAP (As Soon As Possible)
 - ▶ The problem with bounded resources is NP-hard
- ▶ With communication costs:
 - ▶ Both problems are NP-hard
- ▶ Heuristics are required to:
 - ▶ Set priorities to tasks
 - ▶ Allocate tasks to the processors

List heuristics for task graphs

- ▶ Idea: at each instant, allocate as many tasks as possible to the available processors (greedy algorithms)
- ▶ Defining priorities is required if there are more tasks than available processors
- ▶ Generic list scheduling algorithm
 - ▶ 1) Initialization
 - Compute the priority level of all the tasks
 - Set the priority queue as the list of **free** tasks (tasks without predecessors) sorted by decreasing priorities
 - ▶ 2) While it remains tasks to be executed
 - Add the new free tasks to the priority queue
 - If there are **q** available processors and **r** tasks in the queue:
 - Remove the $\min(q, r)$ first tasks of the queue and execute them

Tasks prioritization

- ▶ Prioritization based on the critical path:
 - ▶ Critical path of a task t : weight of the heaviest path starting from t
 - ▶ Idea: prioritize tasks with the heaviest critical path
- ▶ Example DAG:




| Tasks | T_1 | T_2 | T_3 | T_4 | T_5 | T_6 |
|------------|-------|-------|-------|-------|-------|-------|
| Weights | 3 | 2 | 1 | 3 | 4 | 6 |
| Crit. path | 3 | 6 | 7 | 3 | 4 | 6 |

Initial priority queue: (T_3, T_2, T_1)

Communication costs



- ▶ DAG example:
 - ▶ Tasks weights in blue ; communication costs in red



- ▶ Suppose as many processors as needed (4)
 - Sequential execution (no data transfer): 6 tops
 - Parallel execution: 10 tops (critical path: $T_1 \rightarrow T_3 \rightarrow T_4$)
- ▶ A trade-off between, parallelization and communication costs has to be found

Heterogeneous resources

- ▶ Heterogeneous Earliest-Finish-Time (HEFT)
- ▶ List scheduling HEFT
 - ▶ Ordering
 - Set the weights of the tasks
 - Set the weights of the edges
 - Compute the rank (critical path) of each task.
 - Sort the tasks into a list L by non increasing order of their rank
 - ▶ Mapping
 - While the list L of tasks is not empty
 - Select the first task t of the list L
 - Select the resource r that has the earliest finish time for t
 - Allocate task t on resource r
 - Remove t from list L.

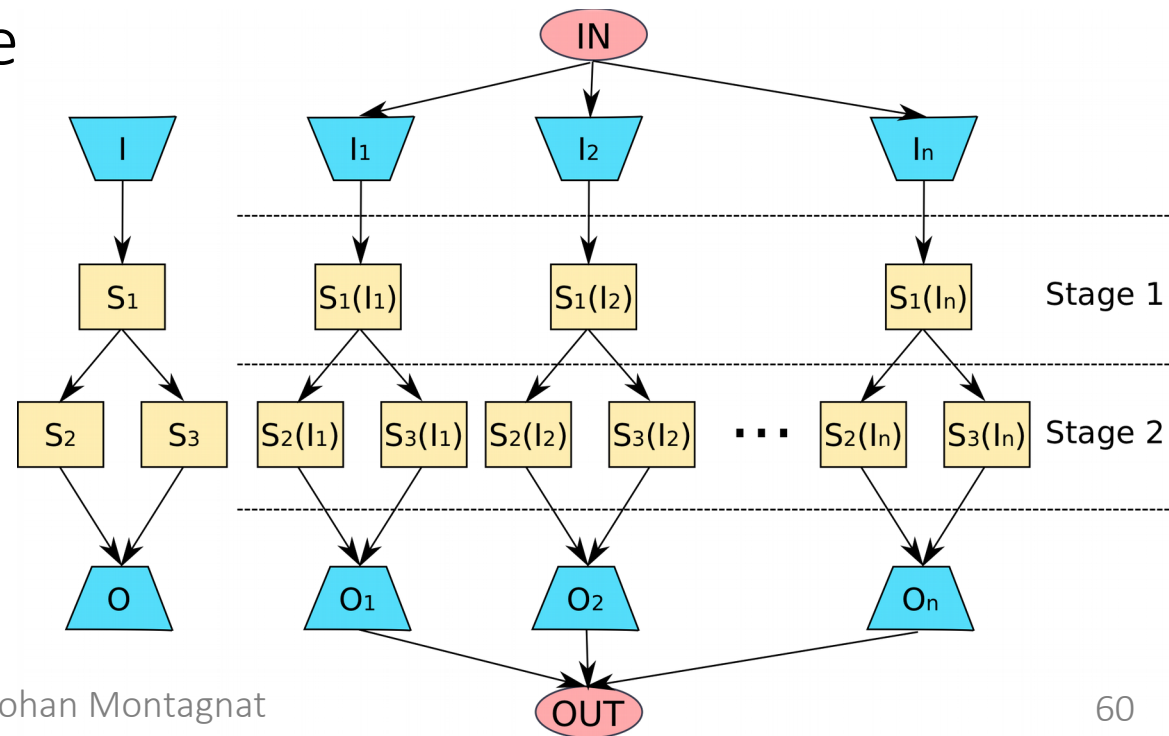


Scientific workflows scheduling

- ▶ Data intensive, scientific workflows
 - ▶ Each workflow activity (service) generates a large number of similar computation tasks: bags of tasks
 - ▶ Static workflow case: a critical path can be estimated
 - ▶ Dynamic workflow case:
 - Some non-deterministic activities (e.g. conditionals) break predictability
 - Sub-workflows without non-deterministic activities can be considered static
- ▶ On large batch distributed systems
 - ▶ Use statistics to estimate tasks execution and data transfer time
- ▶ On clouds
 - ▶ Estimate the number of resources to allocate
 - ▶ May vary during workflow execution time

Workflow-based resource planner for clouds or pilot jobs

- ▶ Executing application in several stages
- ▶ Hypothesis
 - ▶ Known duration of each task (low / predictable variability)
 - ▶ Resources are dedicated to the task (cloud computing)
 - ▶ No loop nor other unbounded control structures (deterministic)
- ▶ Optimizing number of computing resources and network bandwidth for each stage



Workflow cost computation model

► T_i : estimated execution time of stage i

► Computing resources

$$C_r = c_r \times \sum_{i=1}^s m_i \times (Td_i + T_i)$$

► Network bandwidth

$$C_b = c_b \times \sum_{i=1}^s (Td_i + T_i) \sum_{j=1}^{k_i} b_{i,j}$$

► Total execution cost

$$C = C_r + C_b$$

S Number of execution stages

m_i Number of nodes used in stage i ($m_i \leq m_{\max}$)

Td_i Deployment time of stage i

T_i Estimated execution time of stage i (in second)

c_r Per-second cost of a computing node

k_i Number of links used in stage i

$b_{i,j}$ Bandwidth of link j used in stage i (in Mbps)


c_b Per-Mbps cost of bandwidth

Grouping strategy

- ▶ To reduce the impact of jobs submission overhead


- ▶ Activities grouping

- ▶ 6 services – 2 grouped pairs
- ▶ 4 job submissions/input data set




- Recursive grouping

- 4 services – 3 grouped pairs
- 1 job submission/input data set



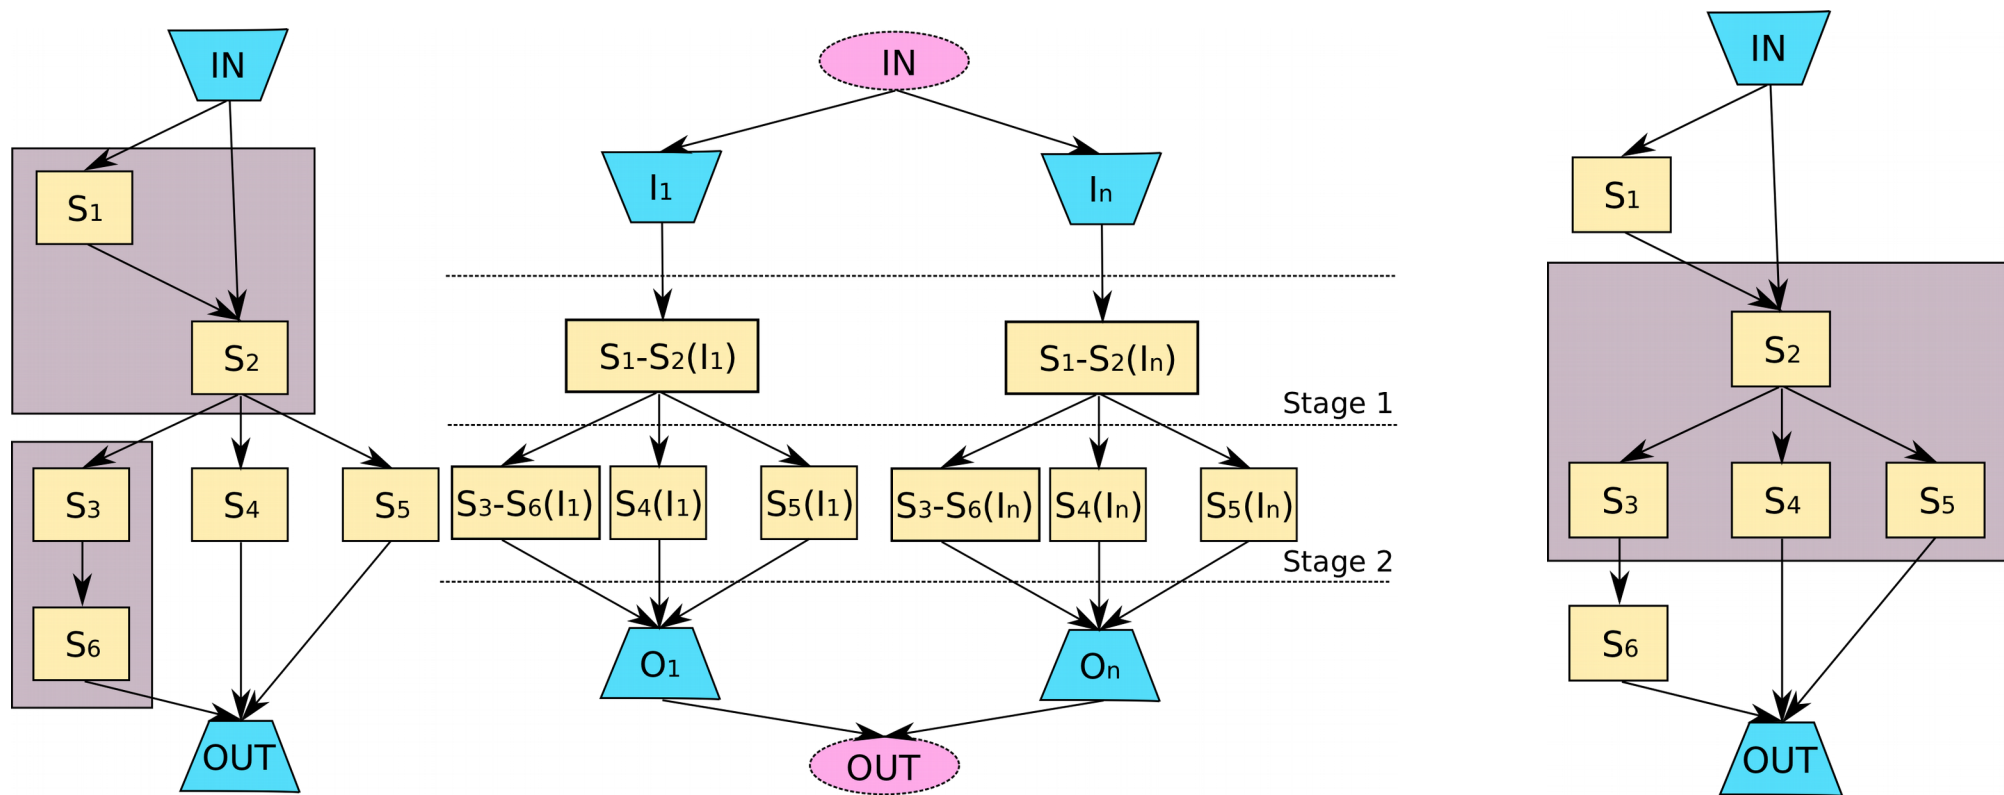
Grouping strategy

- ▶ Let A be a service of the workflow and $\{B_0, \dots, B_n\}$ its children
- ▶ For grouping A and B_{i0} : no parallelism loss \Leftrightarrow (1) & (2)
 - ▶ (1) B_{i0} is an ancestor of every B_j
 - ▶ (2) Every ancestor of B_{i0} is an ancestor of A (or A itself)
- ▶ No parallelism loss \Rightarrow (1) & (2)
 - ▶ $\neg(1) \Rightarrow$ parallelism between B_j and B_{i0} is broken
 - ▶ $\neg(2) \Rightarrow$ parallelism between A and C is broken
- ▶ (1) & (2) \Rightarrow no parallelism loss
- ▶ This rule is recursively applied on the workflow graph



Grouping vs parallelism loss

- ▶ Aggressive grouping leads to parallelism loss
- ▶ Estimate whether grouping gain compensates for parallelism



Resources mapping

Condor matchmaker example

- ▶ Workload management
 - ▶ Heterogeneous resources
- ▶ Deliver High Throughput Computing
 - ▶ For many experimental sciences, the computing throughput matters. Focus is not instantaneous computing power but the amount of computing that can be harnessed over a long period.
 - ▶ HTC is a 24-7-365 activity: fault tolerance is critical
- ▶ Batch-oriented system
 - ▶ Batch extended with Job Control Languages to face grid heterogeneity
- ▶ Distributed computing IS difficult
 - ▶ team of ~35 faculty, full time staff and students (U. Wisconsin)
 - ▶ established in 1985
 - ▶ Faces software/middleware engineering challenges in a UNIX/Linux/Windows/OS X environment




Matchmaking heterogeneous resources

- ▶ Run jobs in a variety of environments
 - ▶ Local dedicated clusters (machine rooms)
 - ▶ Local opportunistic (desktop) computers
 - ▶ Grid environments; Interface to other systems
- ▶ Matchmaking process

I need a mac for this code
to run


I need a linux box with 2Gb
RAM



Desktop Computers



Dedicated Clusters



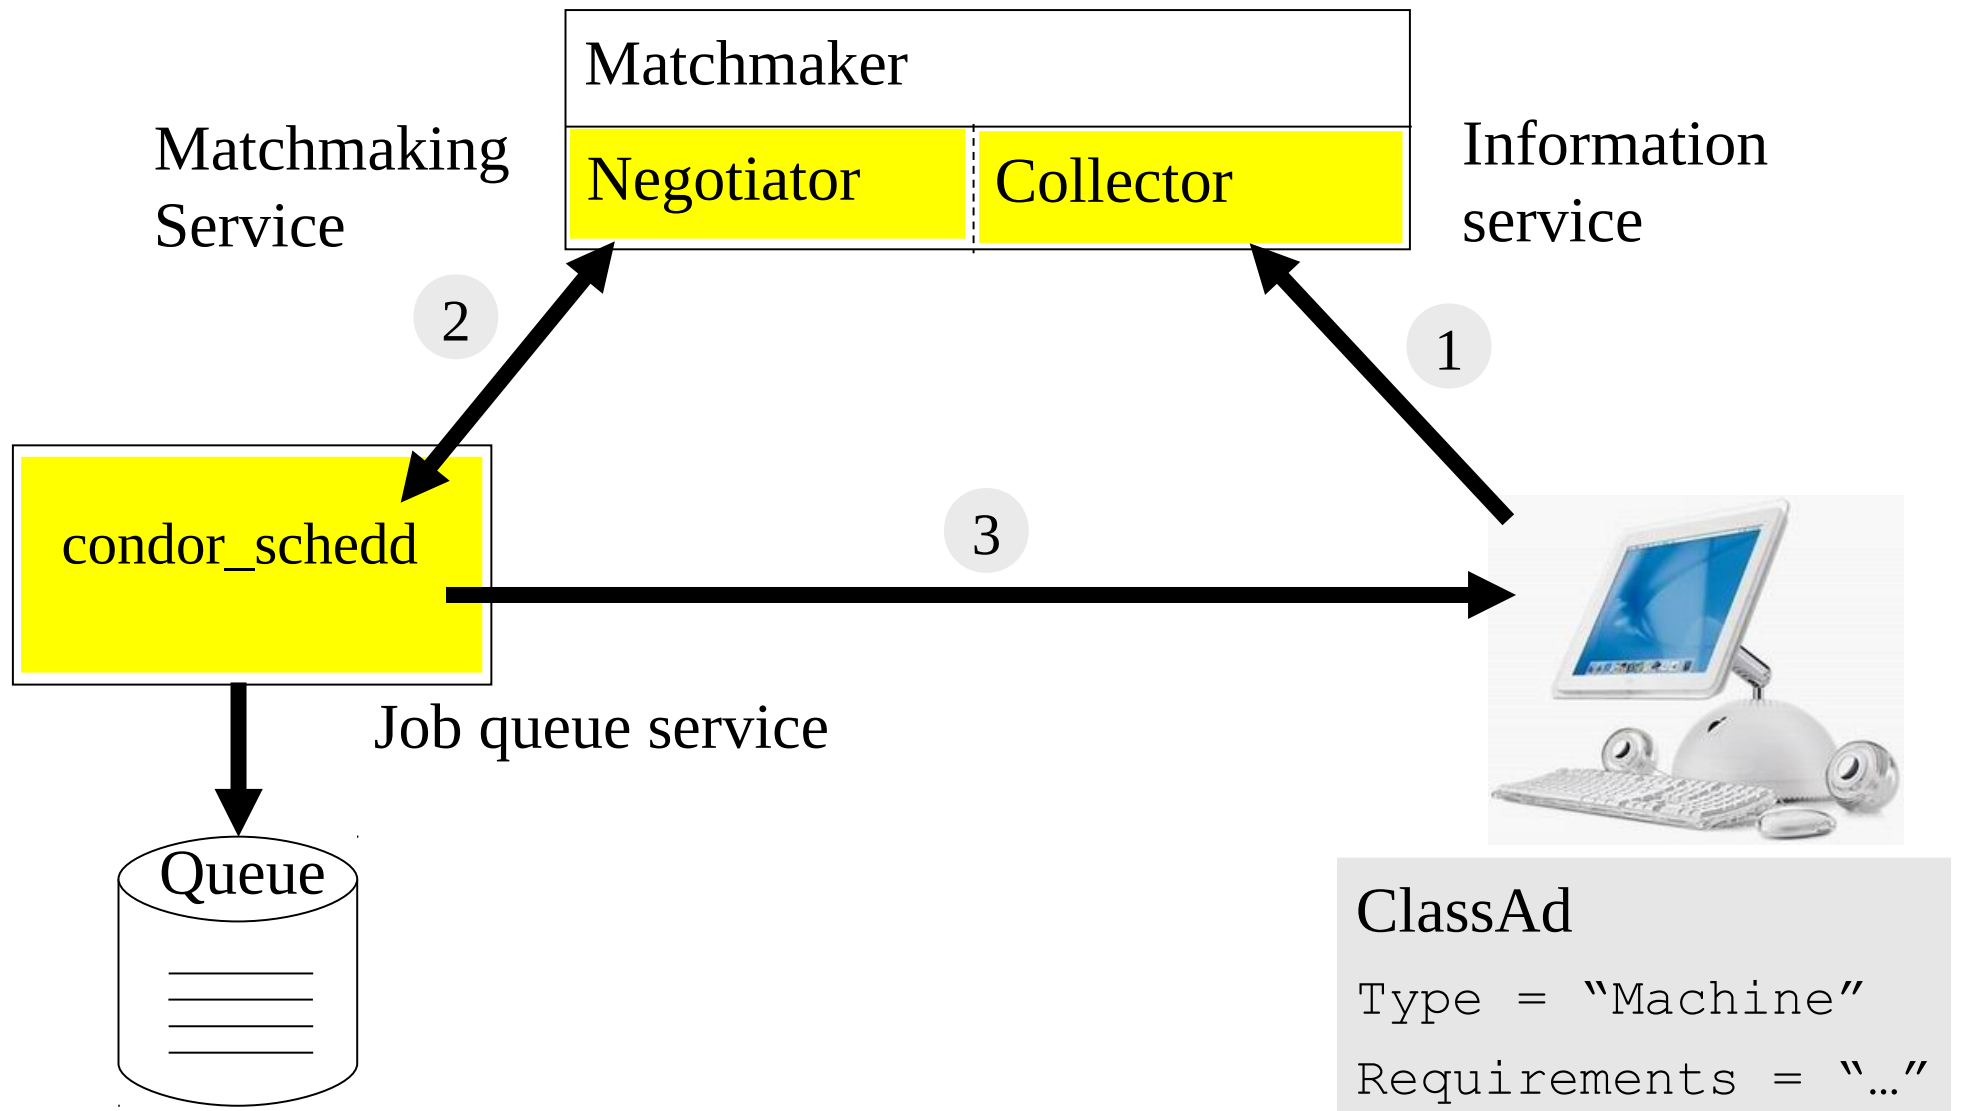
Matchmaking

- ▶ Condor conceptually divides people into three groups
 - ▶ Job submitters
 - I need Linux, and I prefer faster machines
 - ▶ Machine owners
 - I prefer jobs from the physics group
 - I will only run jobs between 8pm and 4am
 - I will only run certain types of jobs
 - Jobs can be preempted if something better comes along
 - ▶ Cluster administrator
 - When can jobs preempt other jobs?
 - Which users have higher priority?
 - Do some groups of users have allocations of computers?

Matchmaking

- ▶ Matchmaking is two-way
 - ▶ Job describes what it requires:
 - I need Linux && 8 GB of RAM
 - ▶ Machine describes what it provides:
 - I will only run jobs from the Physics department
- ▶ Matchmaking allows preferences
 - ▶ I **need** Linux, and I **prefer** machines with more memory but will run on any machine you provide me
- ▶ ClassAds Job Description Language (JDL)
 - ▶ Stating facts
 - Job's executable is analysis.exe
 - Machine's load average is 5.6
 - ▶ Stating preferences
 - I require a computer with Linux

Matchmaking diagram



Including dynamic information (load...)

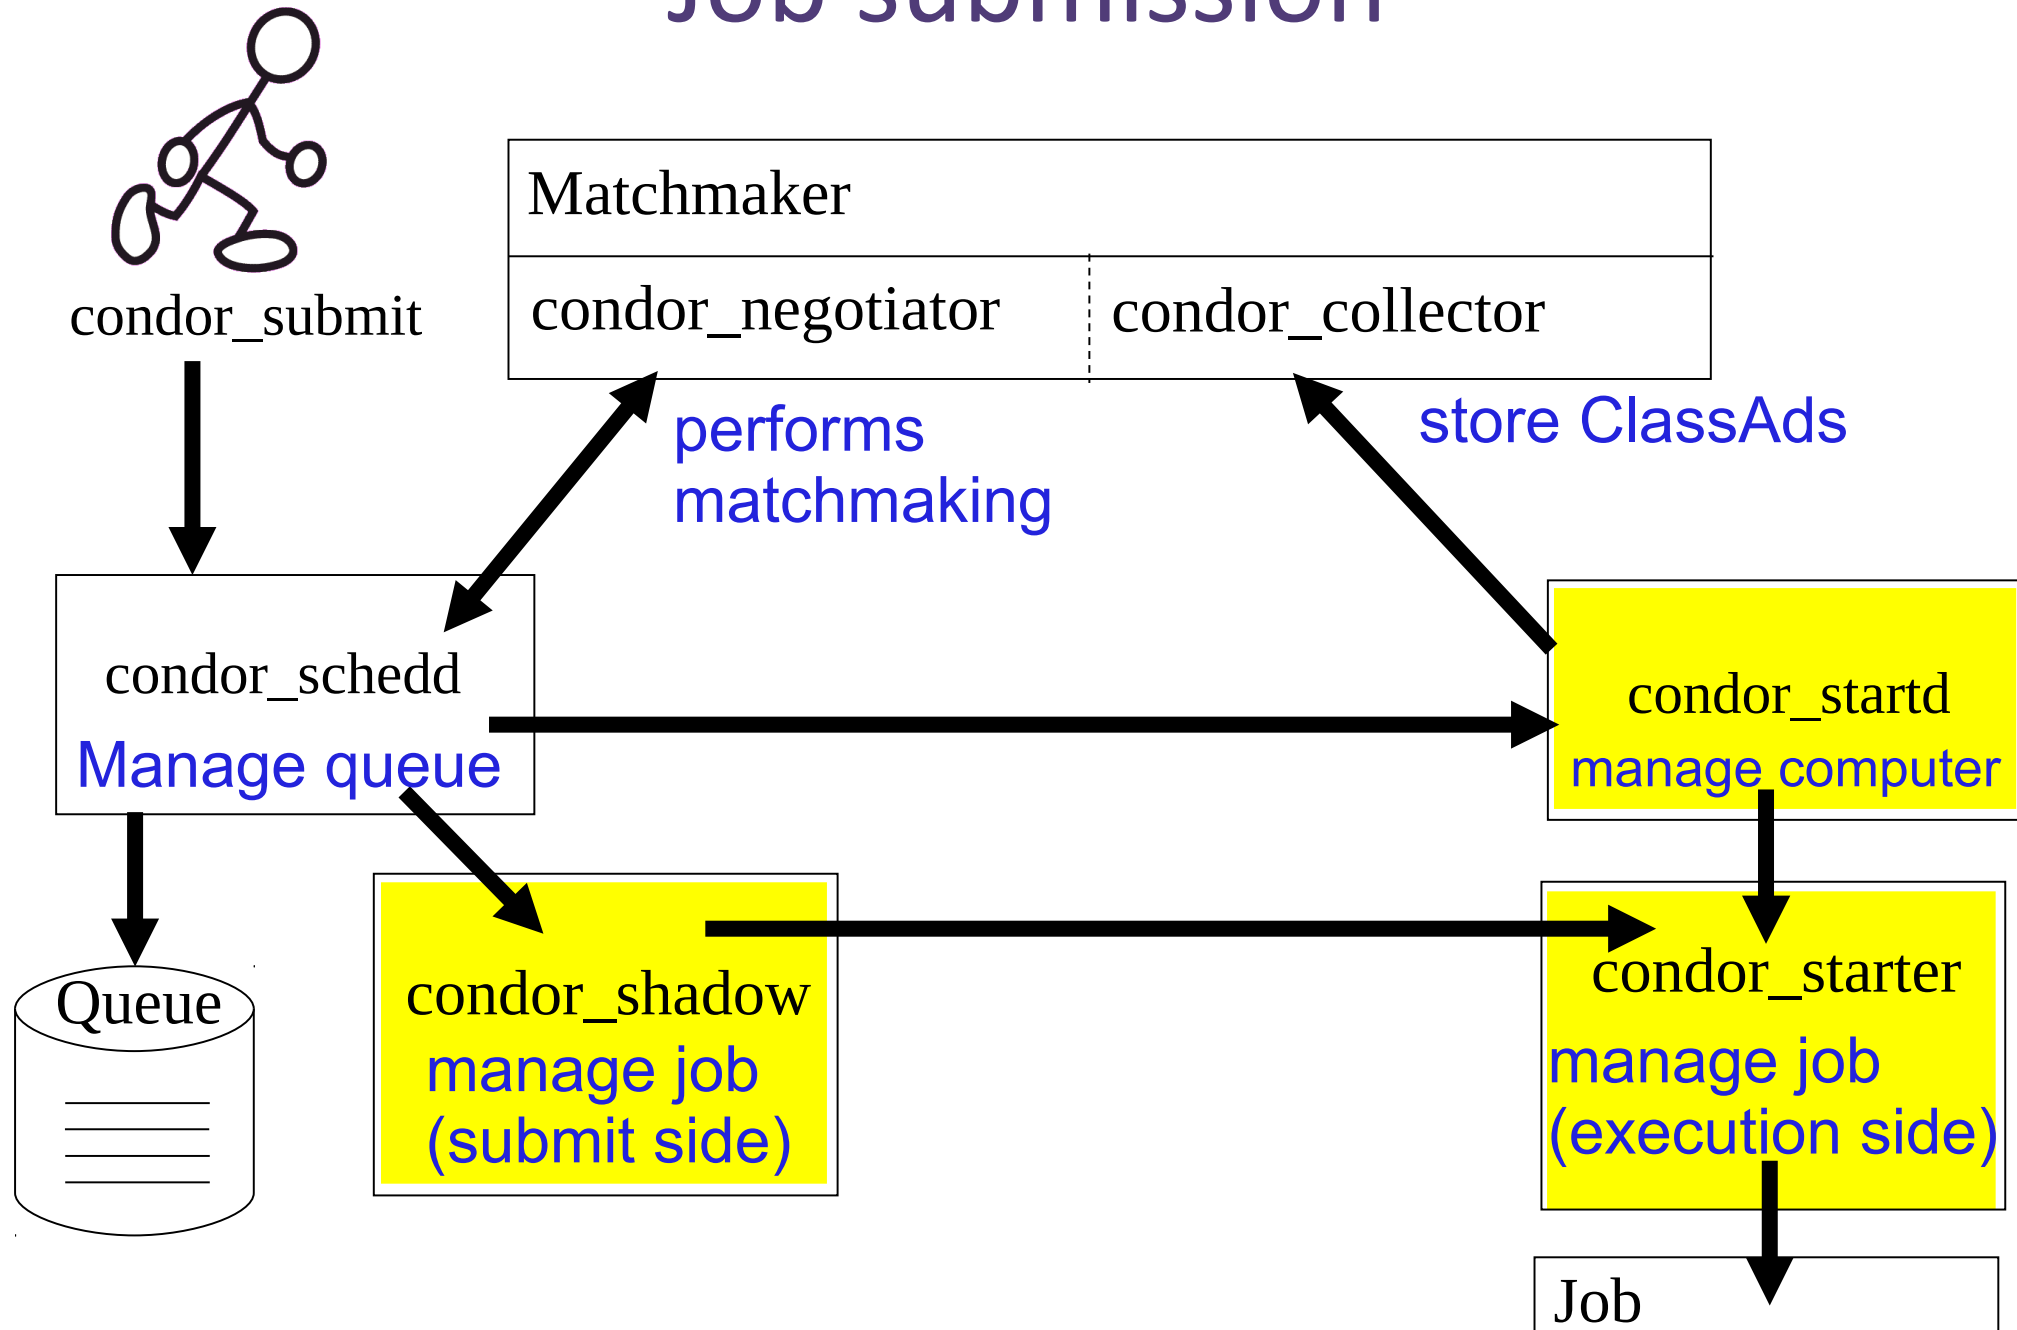
ClassAds JDL

- ▶ ClassAds are:
 - ▶ Semi-structured
 - ▶ Attribute = Expression
 - ▶ Schema-free, user-extensible
- ▶ Extensible declaration
 - ▶ HasJava_1_4 = TRUE
 - ▶ ShoeLength = 7
- ▶ Extensible matchmaking
 - ▶ Requirements =
OpSys == "LINUX" &&
HasJava_1_4 == TRUE

▶ Example

```
MyType           = "Job"
TargetType       = "Machine"
ClusterId        = 1377
Owner            = "roy"
Cmd              = "analysis.exe"
Requirements     =
    (Arch == "INTEL")
    && (OpSys == "LINUX")
    && (Disk >= DiskUsage)
    && ((Memory * 1024) >= ImageSize)
...
```

Job submission



Job submission


- ▶ Users submit jobs to scheduler
 - ▶ Jobs described as ClassAds
 - ▶ Each scheduler has a queue
 - ▶ Scheduler / queues are not centralized
- ▶ Negotiator
 - ▶ collects list of computers
 - ▶ contacts each schedd (What jobs do you have to run?)
 - ▶ compares each job to each computer to find a match
 - Evaluate requirements of job & machine in context of both ClassAds
 - If both evaluate to true, there is a match
- ▶ Fault tolerance scheduler
 - ▶ Resubmission
 - ▶ Fail-over scheduler

Large-scale computing techniques


- ▶ Flocking
 - ▶ Metascheduling: connect scheduler to several condor pools
- ▶ Grid interface
- ▶ Pilot jobs
 - ▶ Job-based reservation of resources and application level scheduling

Flocking

- ▶ Submit a scheduler to several pools
 - ▶ Share condor pools between institutions
 - ▶ Try to run on local pool first, then try to run on remote pool



- ▶ Networking issues with private networks
 - ▶ A communication broker may be needed if the scheduler is not able to communicate directly with every execute node

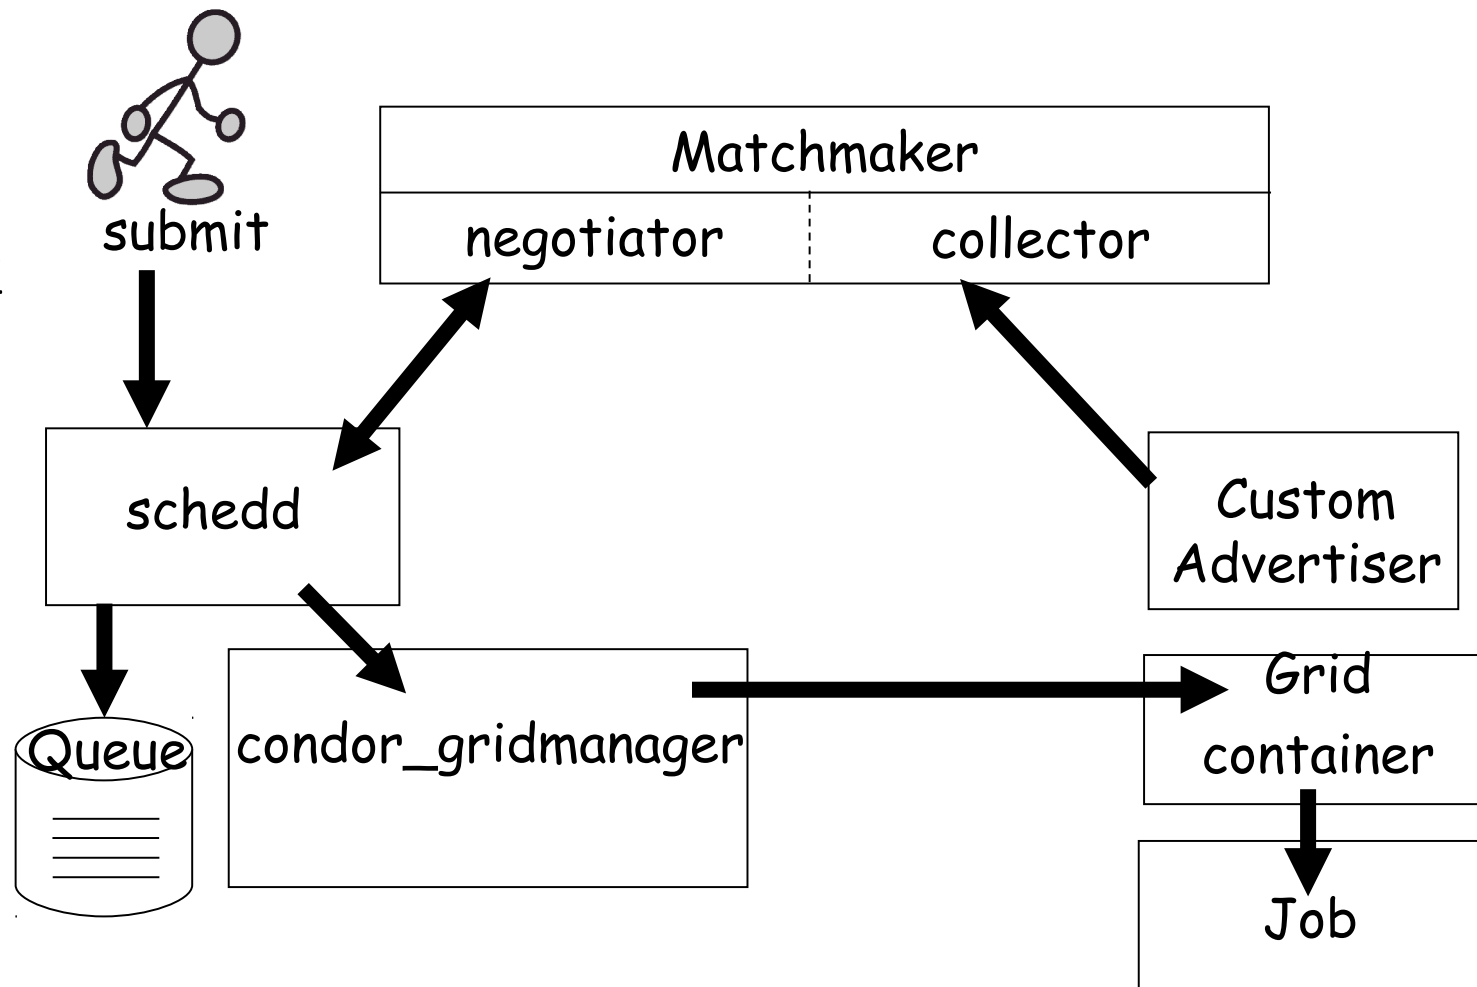


Condor-G

- ▶ Submit jobs to other grid systems
 - ▶ Minimal changes to job description

- ▶ Grids

- ▶ Globus 2
- ▶ Globus 4
- ▶ Amazon EC2
- ▶ Nordugrid
- ▶ Unicore
- ▶ PBS
- ▶ LSF
- ▶ Condor (!)
- ▶ ...

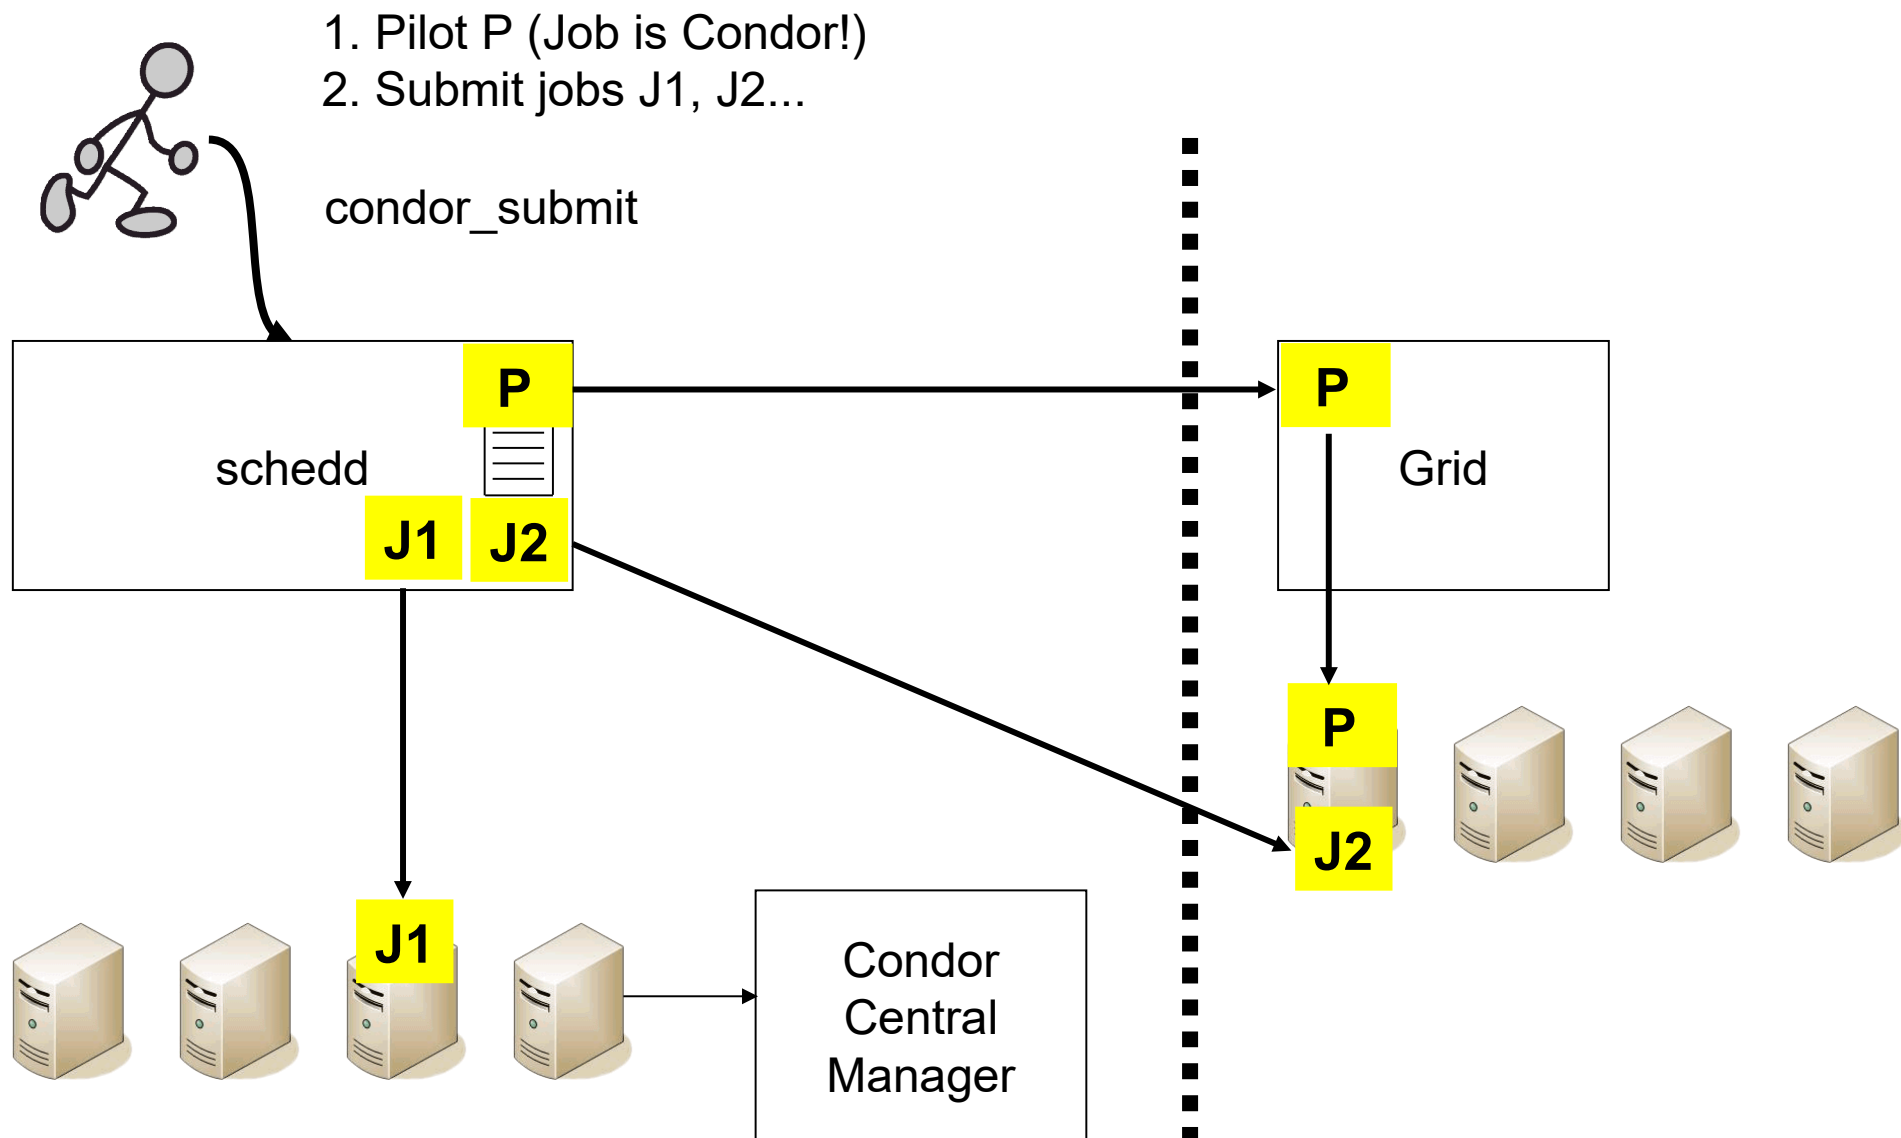


Pilot jobs

- ▶ Pilot jobs are application-level scheduler jobs that once executing on a grid resource schedule other jobs on it
 - ▶ Resources allocation through a custom batch system (bypasses grid workload manager)
 - ▶ Enables application-level scheduling
- ▶ Example: overlaying Condor on another system
 - ▶ Submit startd as a grid job to start a new pilot
 - ▶ Grow the condor pool with the new startd daemon
- ▶ Limitations
 - ▶ The scheduler has to be able to open communication with the pilot jobs
 - ▶ Security is tricky (whose job is ran by the pilot?)
 - ▶ System administrators do not like pilots so much

Condor pilot jobs

- Startd can be run as a grid pilot job (Condor Glide-in)



From Pilots to Clouds

- ▶ Clouds: resources allocation
- ▶ Pilots: resource dedication through job submission

